

SES

Chapter 2: General-Purpose Processors

Prof. Dr.-Ing. Bernd-Christian Renner



Contents

1. Basic Architecture
2. Memory Architectures
3. Application Development
4. Programming
5. Running a Program



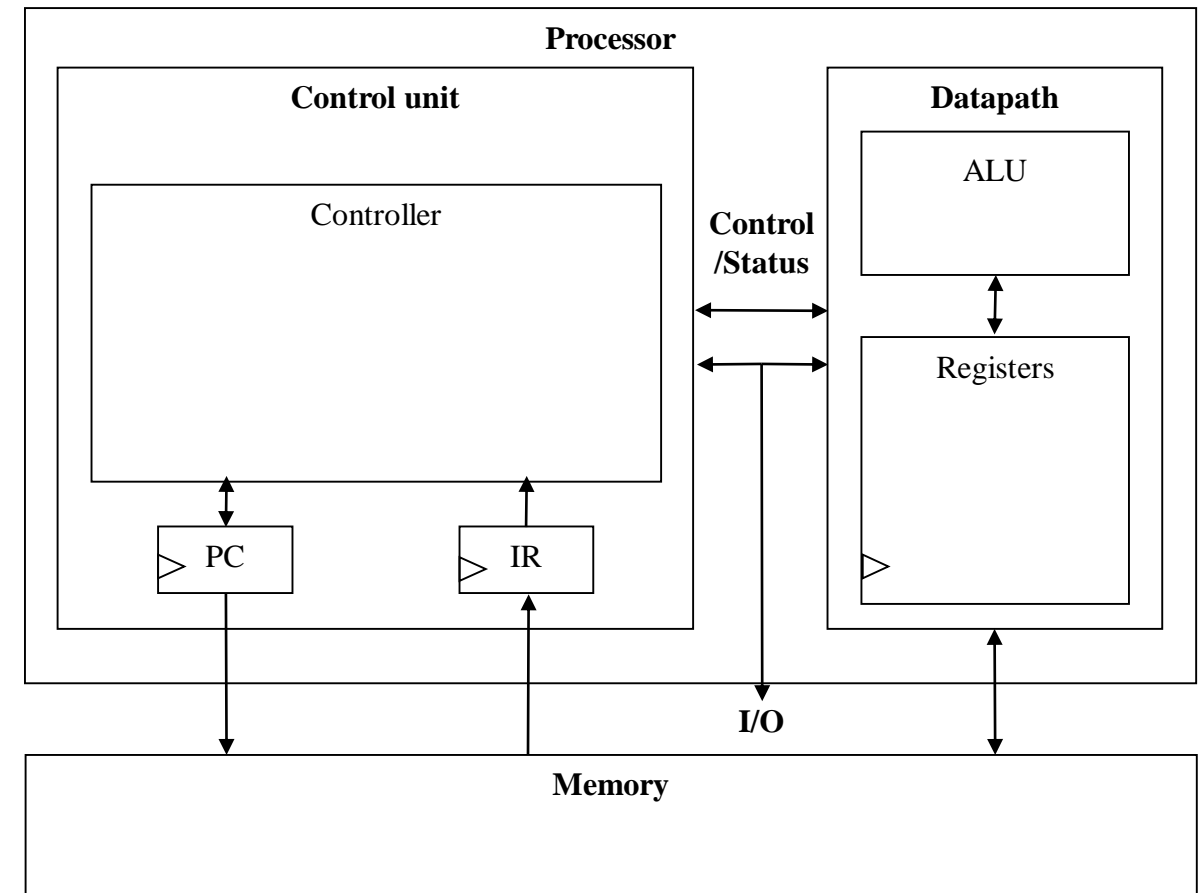
Basic Architecture

Introduction

- General-Purpose Processor
 - Processor designed for a variety of computation tasks (a.k.a. CPU or microprocessor)
 - Low unit cost, in part because manufacturer spreads NRE over large numbers of units
 - Carefully designed since higher NRE is acceptable
 - Can yield good performance, size and power
 - Advantage for embedded systems: Low NRE cost, short time-to-market/prototype, high flexibility
 - User just writes software; no processor design

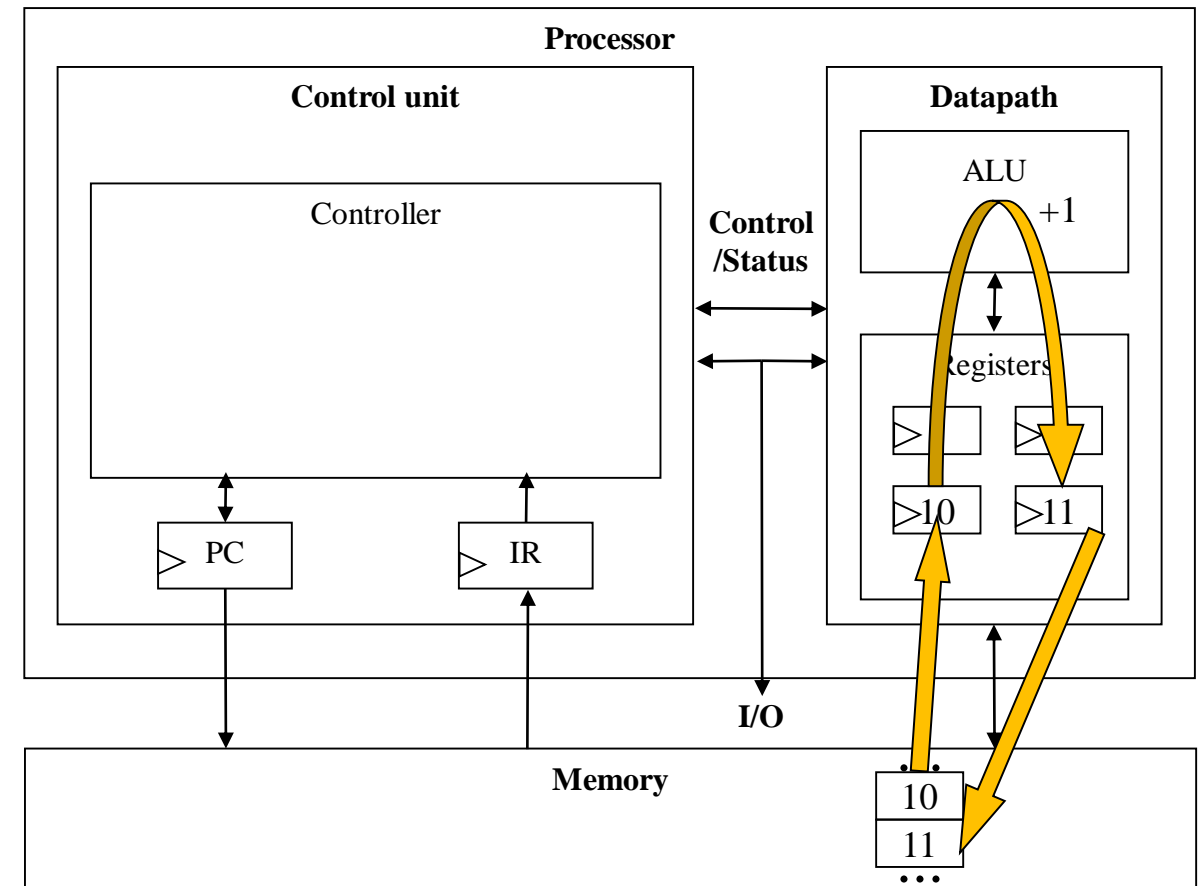
Basic Architecture

- Control unit and datapath
 - Compare with single-purpose processor
- Key differences
 - Datapath is general
 - Control unit doesn't store algorithm
 - algorithm is “programmed” into memory



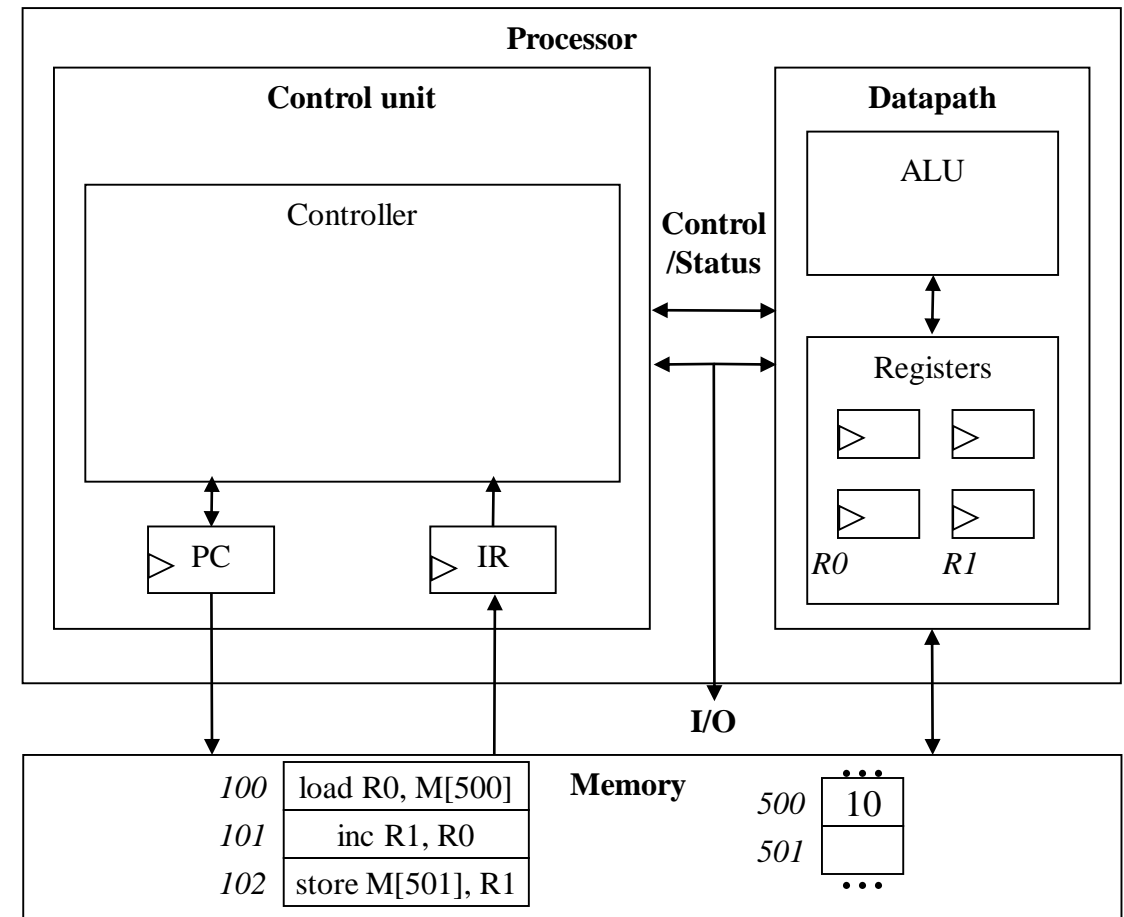
Datapath Operations

- **Load**
 - Read memory location into register
- **ALU operation**
 - Input certain registers through ALU, store back in register
- **Store**
 - Write register to memory location



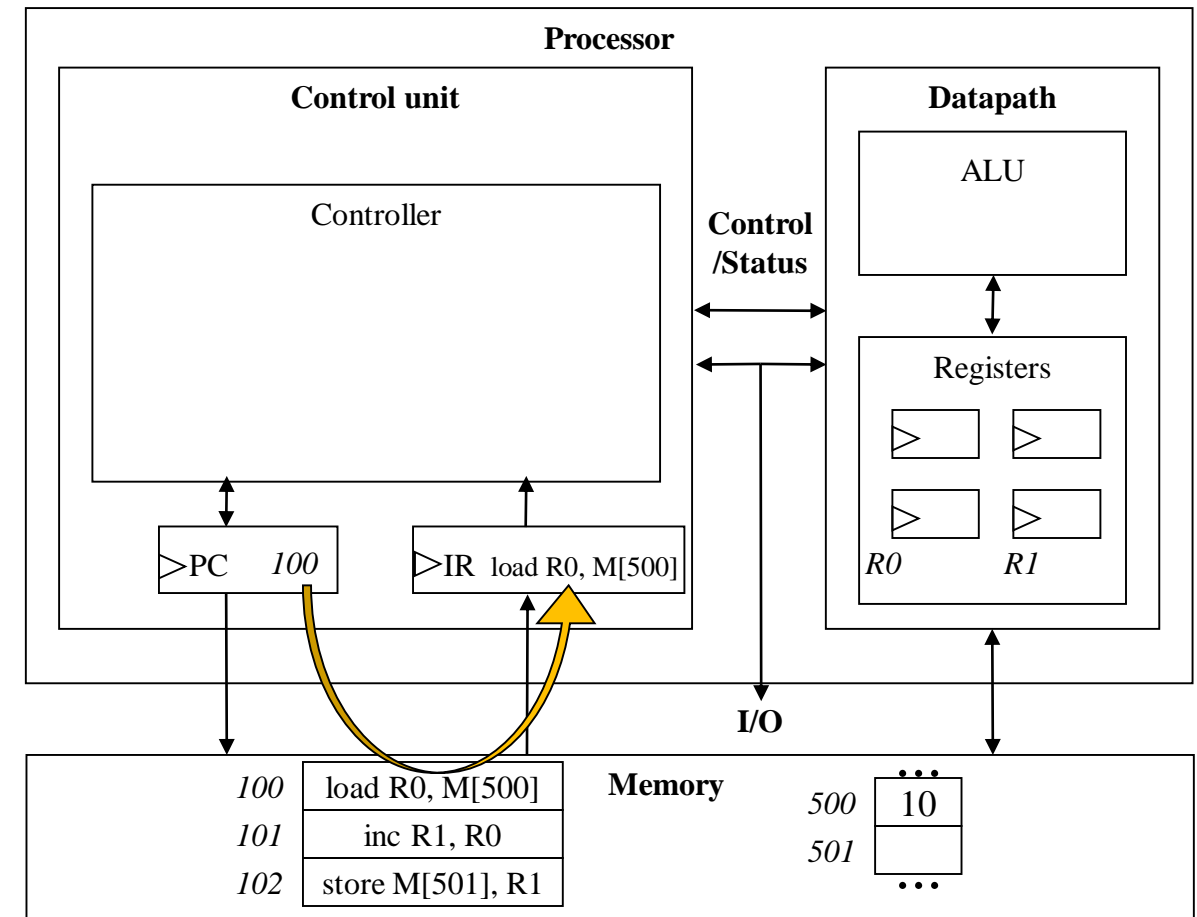
Control Unit

- Control unit: configures datapath operations
 - Sequence of desired operations (“instructions”) stored in memory – “program”
- Instruction cycle – broken into several sub-operations, each one clock cycle, e.g.:
 - Fetch: Get next instruction into IR
 - Decode: Determine what the instruction means
 - Fetch operands: Move data from memory to datapath register
 - Execute: Move data through ALU
 - Store results: Write data from register to memory



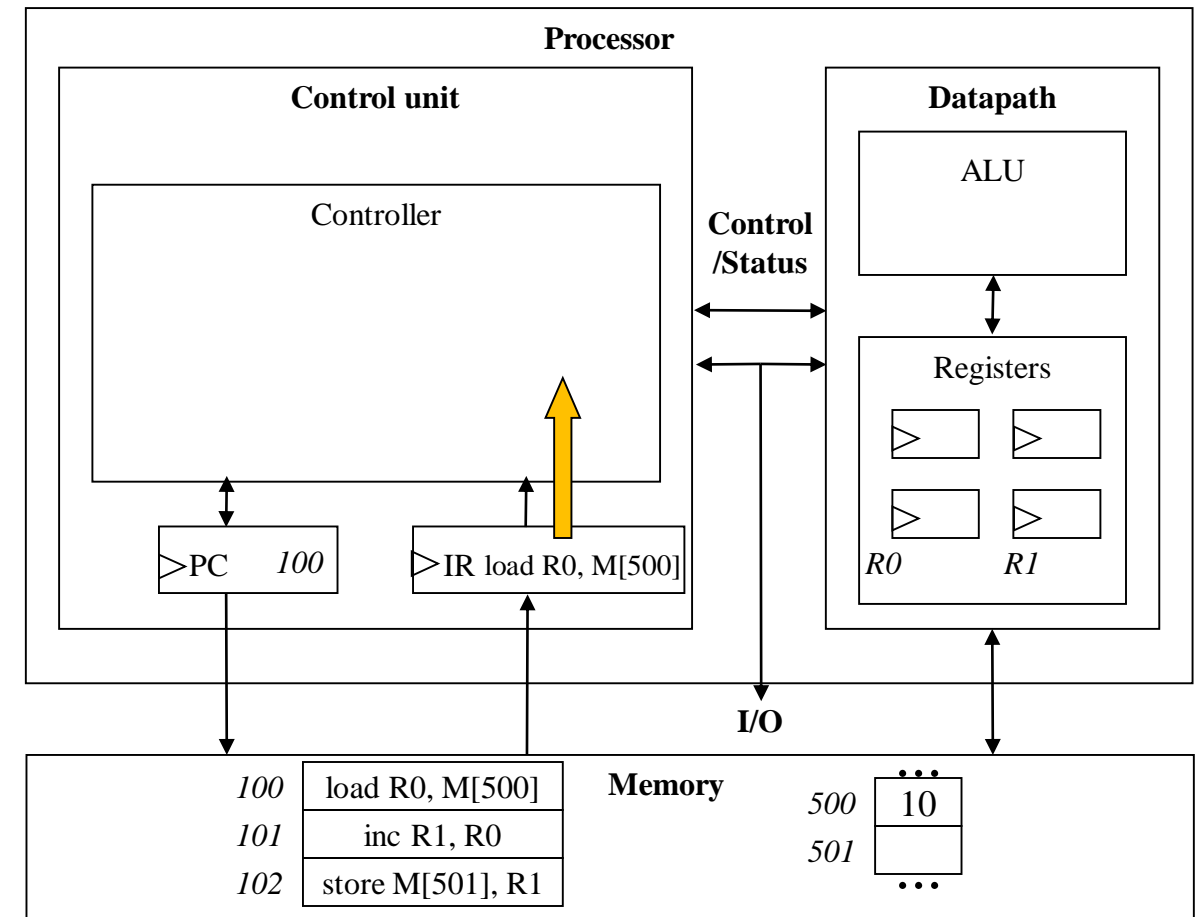
Control Unit Sub-Operations

- Fetch
 - Get next instruction into IR
 - PC: program counter, always points to next instruction
 - IR: holds the fetched instruction



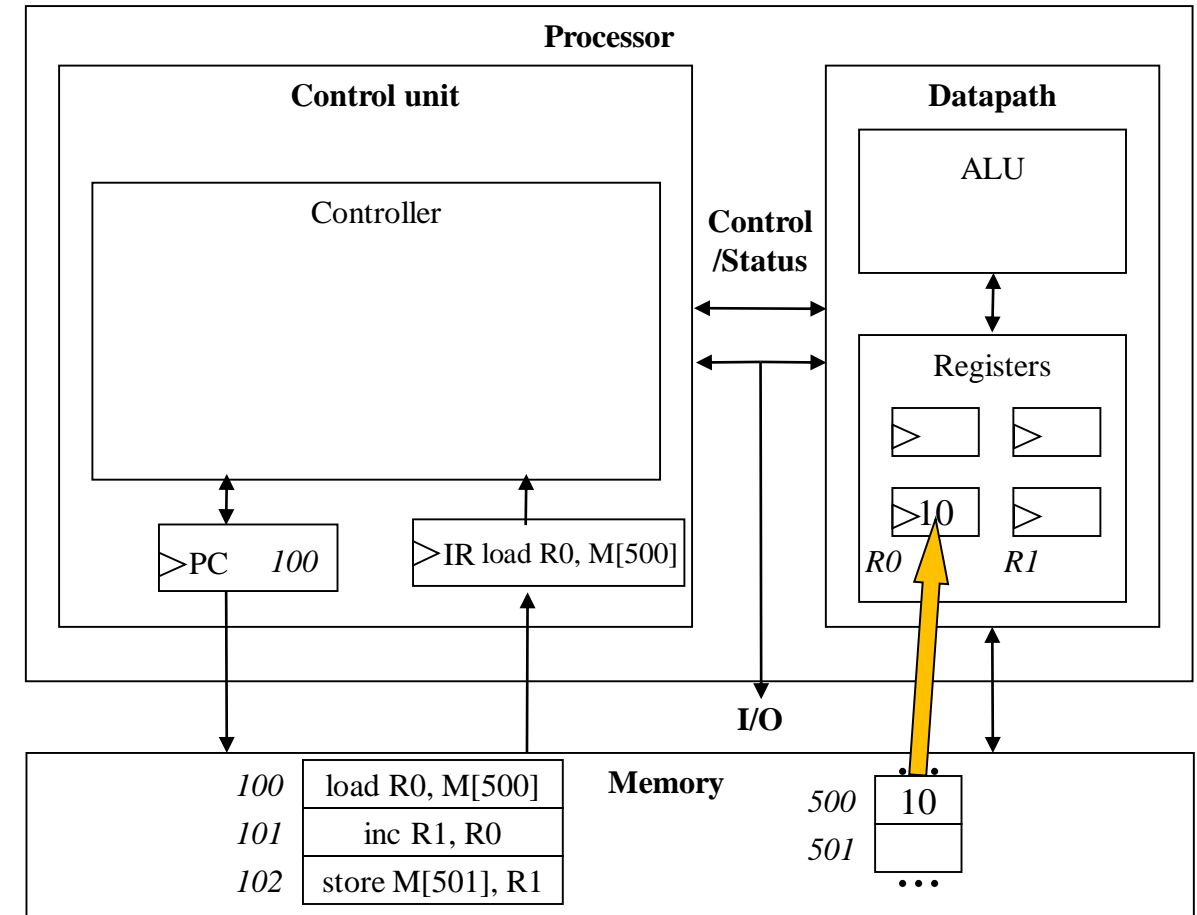
Control Unit Sub-Operations

- Decode
 - Determine what the instruction means



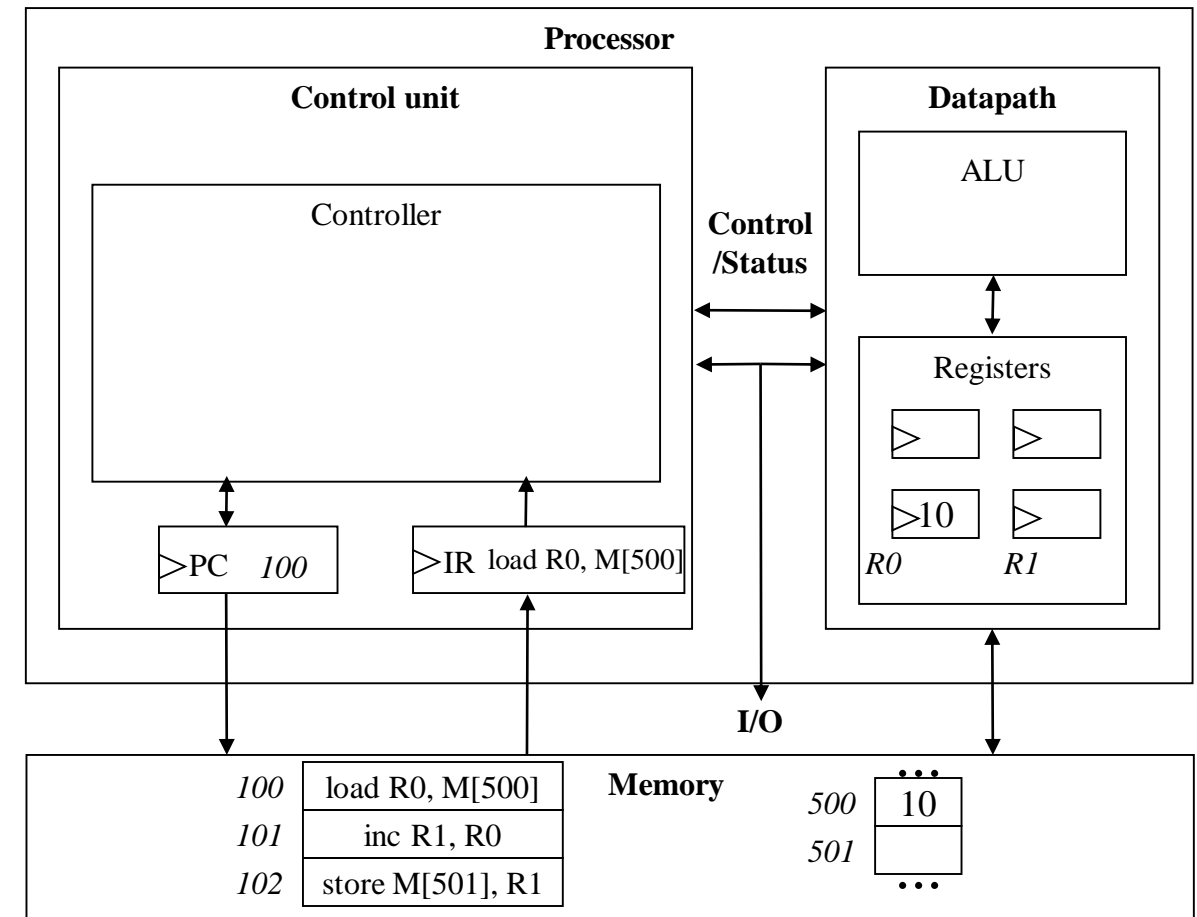
Control Unit Sub-Operations

- Fetch operands
 - Move data from memory to datapath register



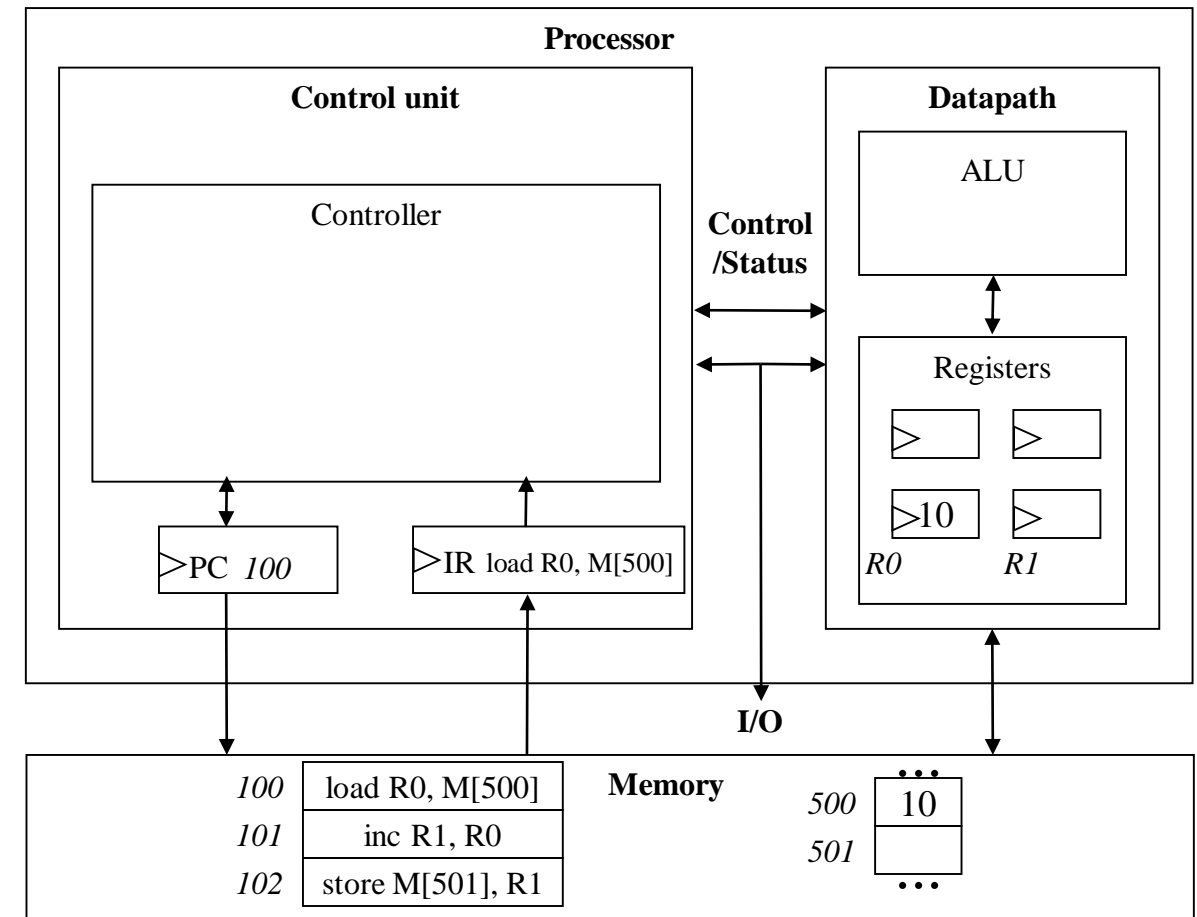
Control Unit Sub-Operations

- Execute
 - Move data through the ALU
 - This particular instruction does nothing during this sub-operation

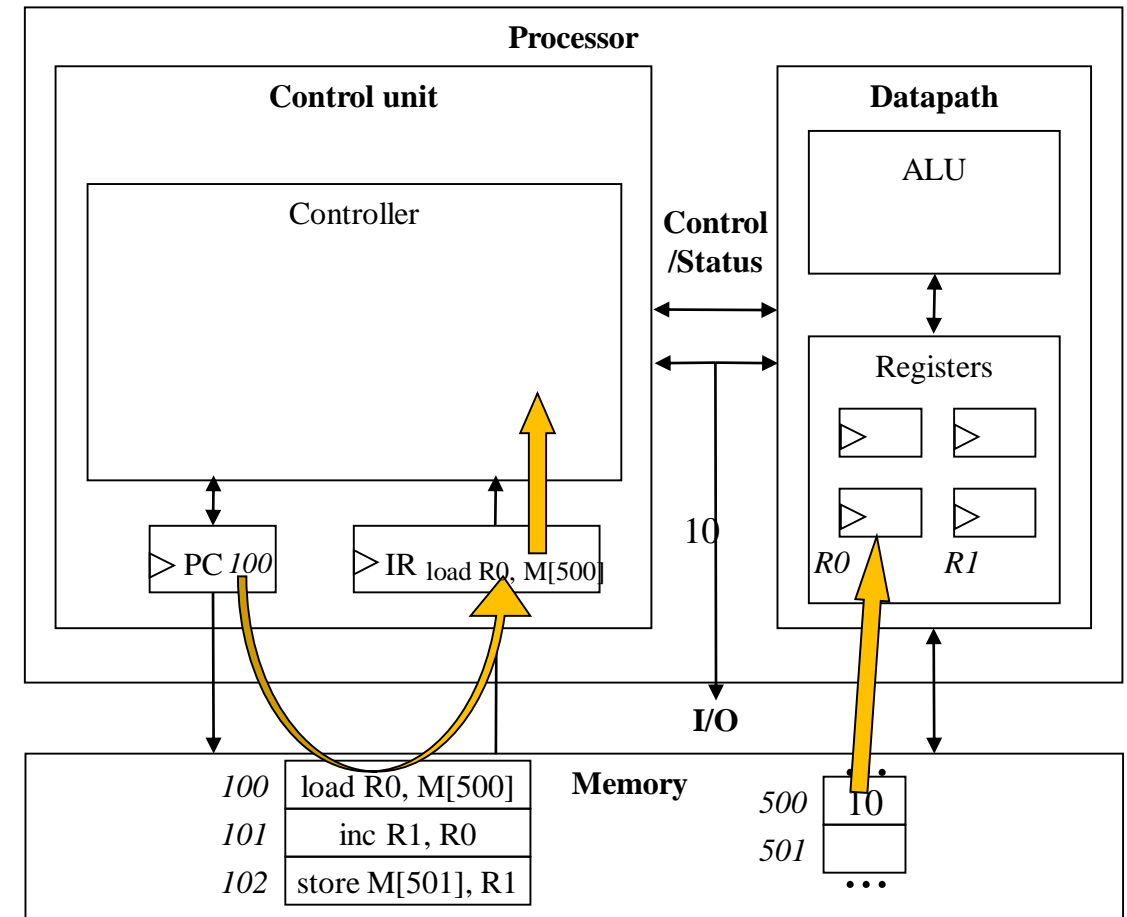



Control Unit Sub-Operations

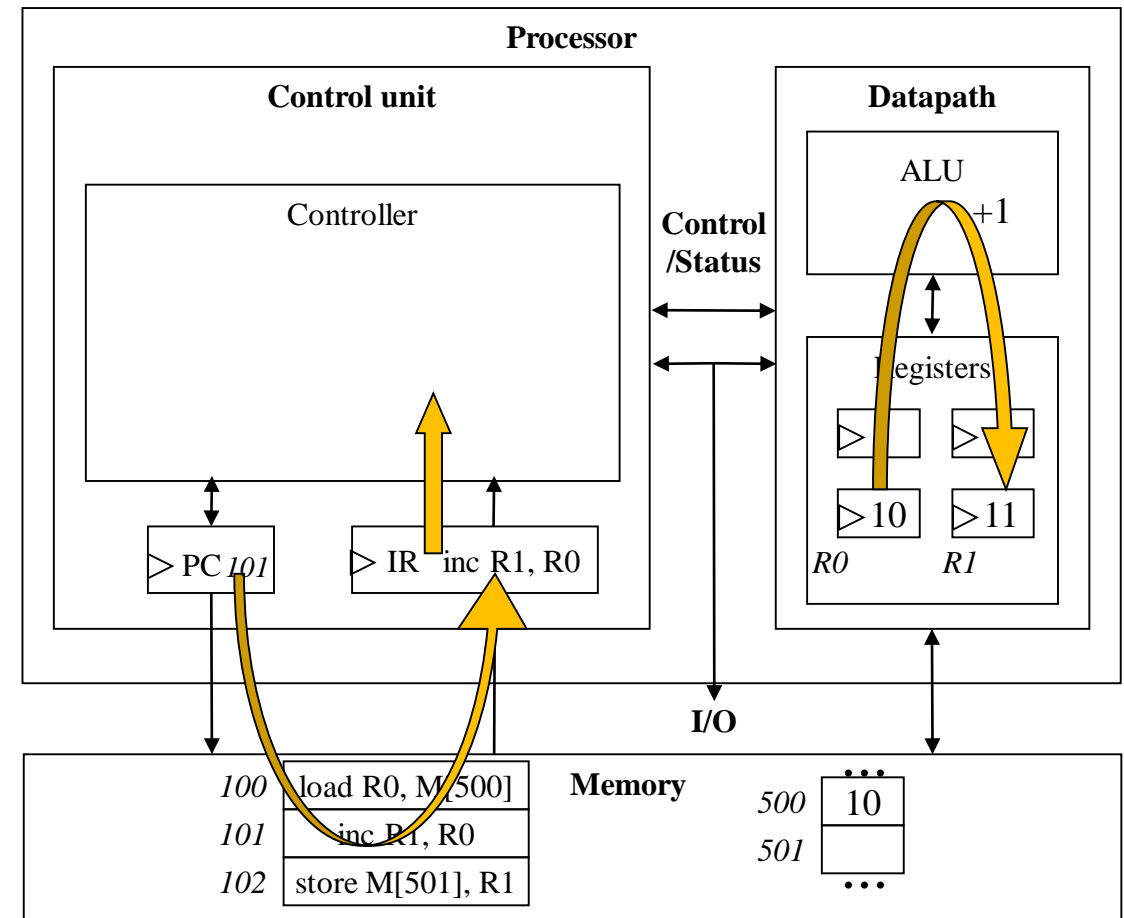

- Store results
 - Write data from register to memory
 - This particular instruction does nothing during this sub-operation



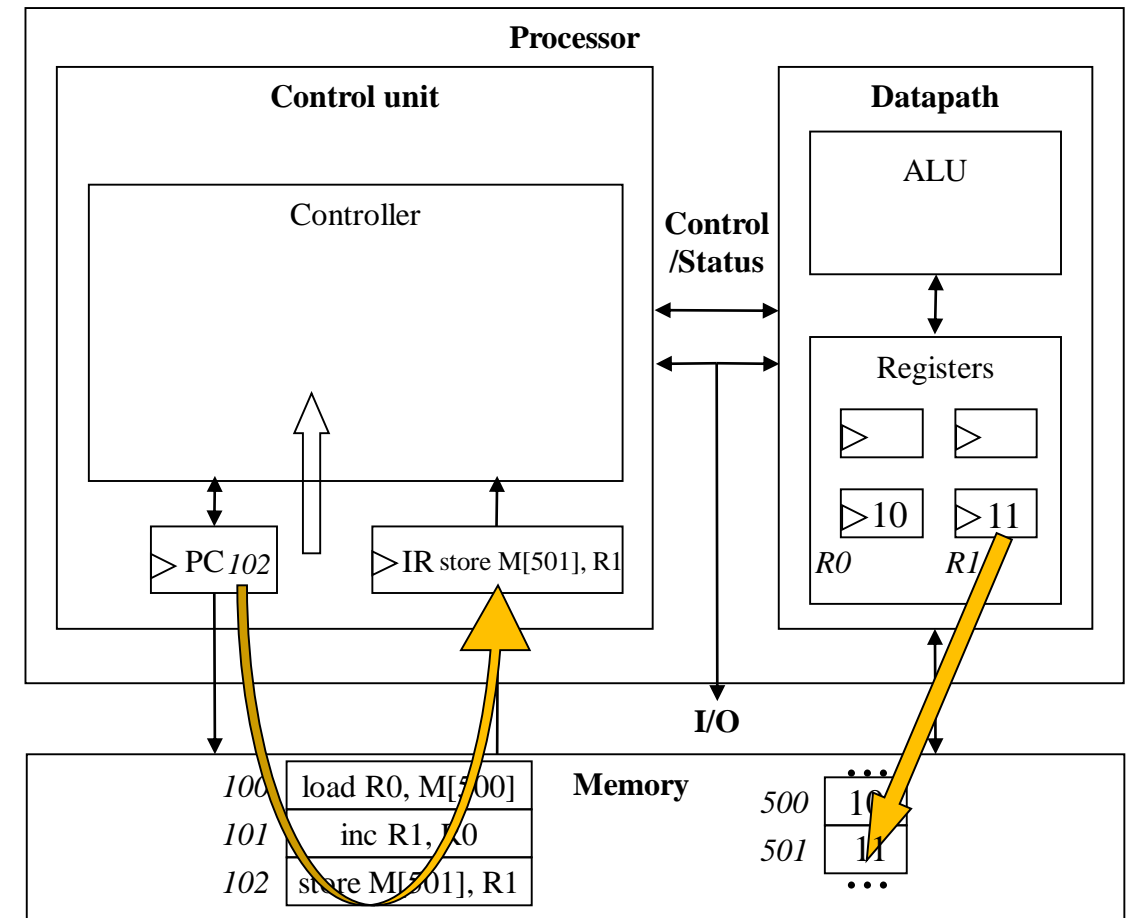
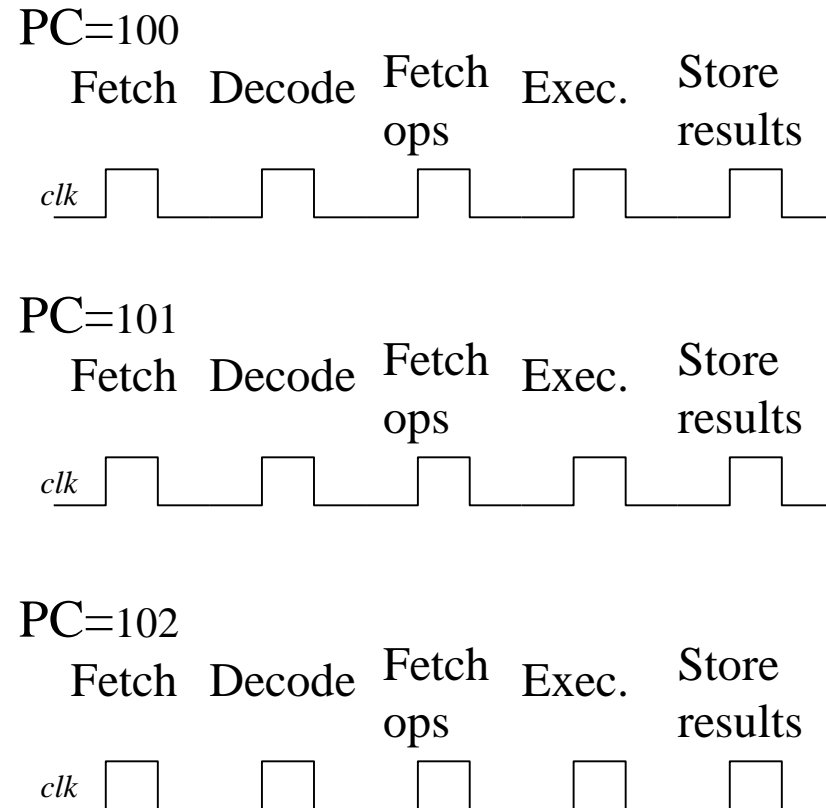
Instruction Cycles



Instruction Cycles

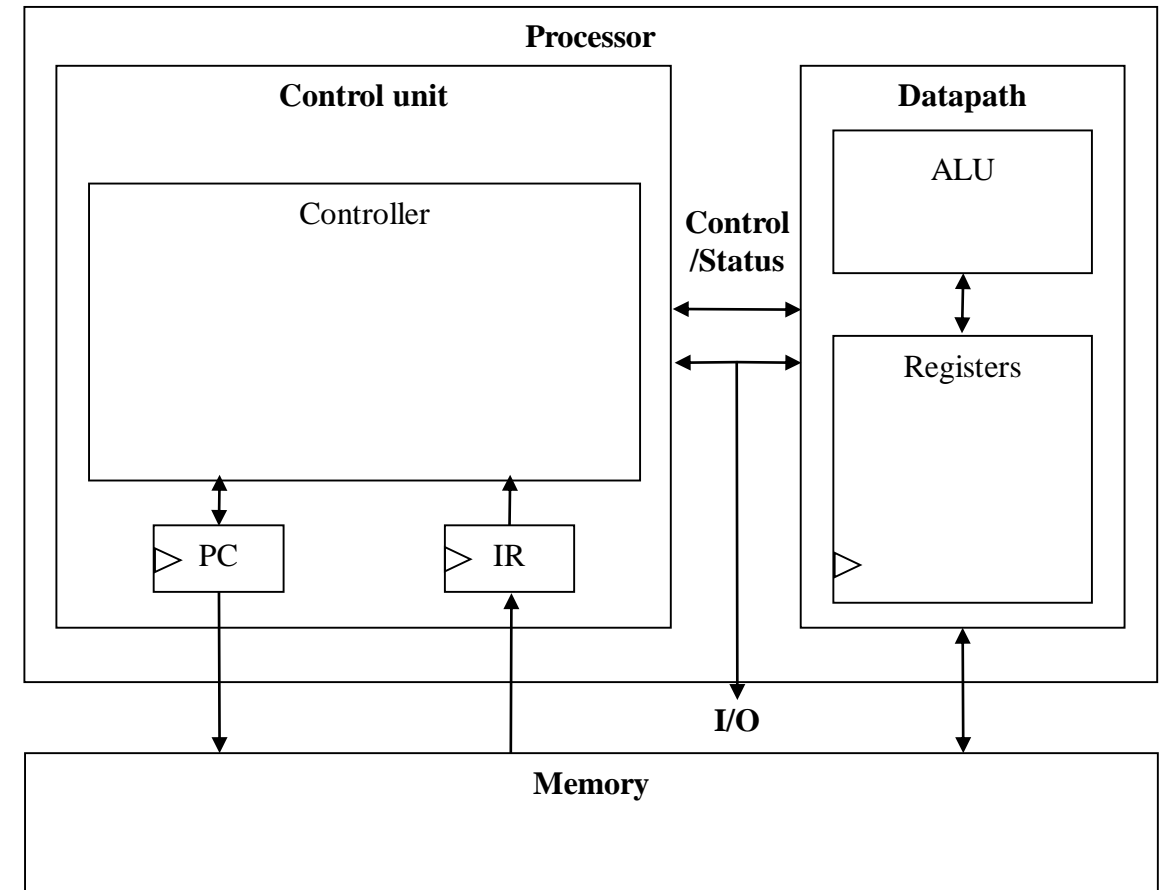


Instruction Cycles



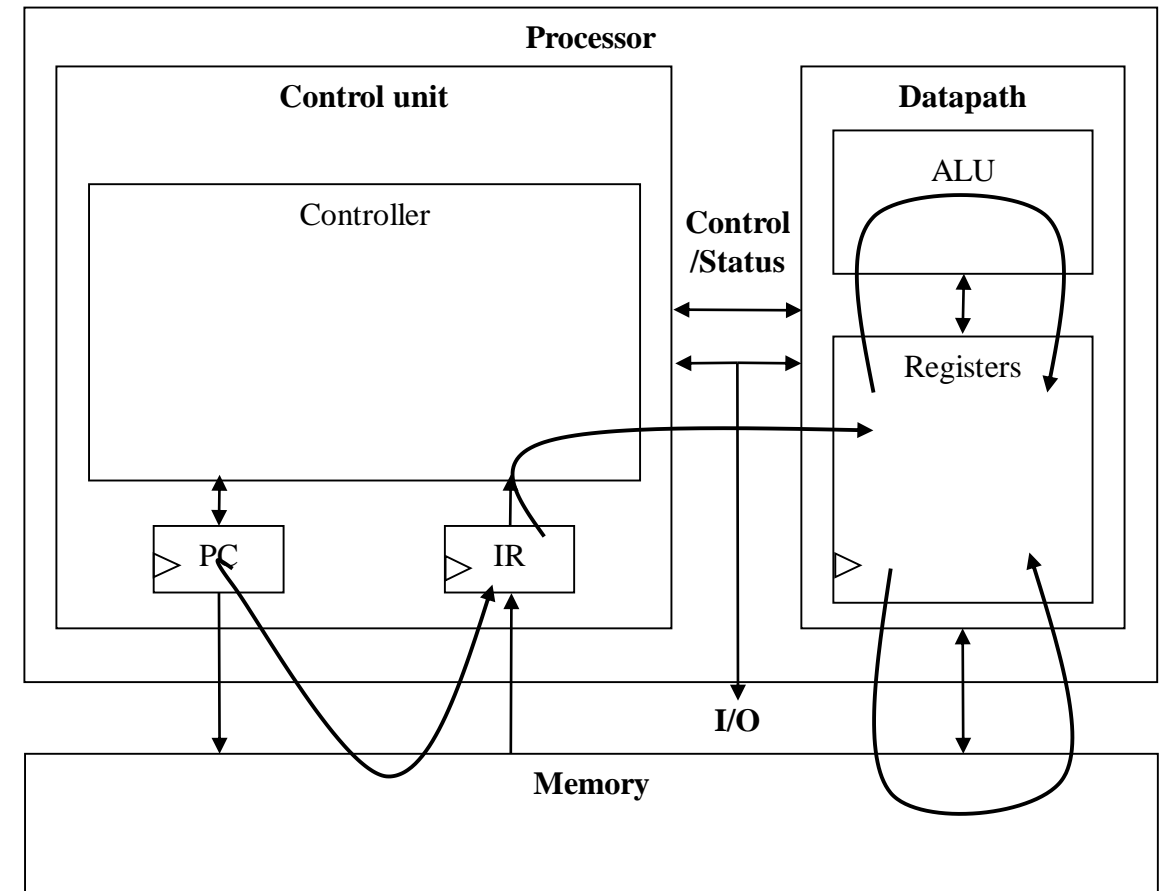
Architectural Considerations

- N-bit processor
 - N-bit ALU, registers, buses, memory data interface
 - Embedded: 8-bit, 16-bit, 32-bit common
 - Desktop/servers: 32 or 64 bit
- PC size determines address space
 - Often larger than data word size
 - PC size M, then address space is 2^M

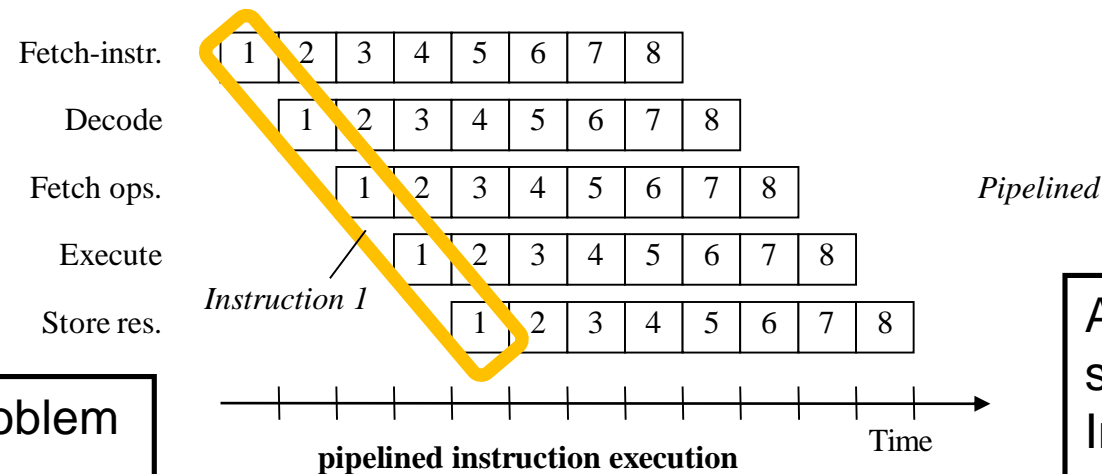
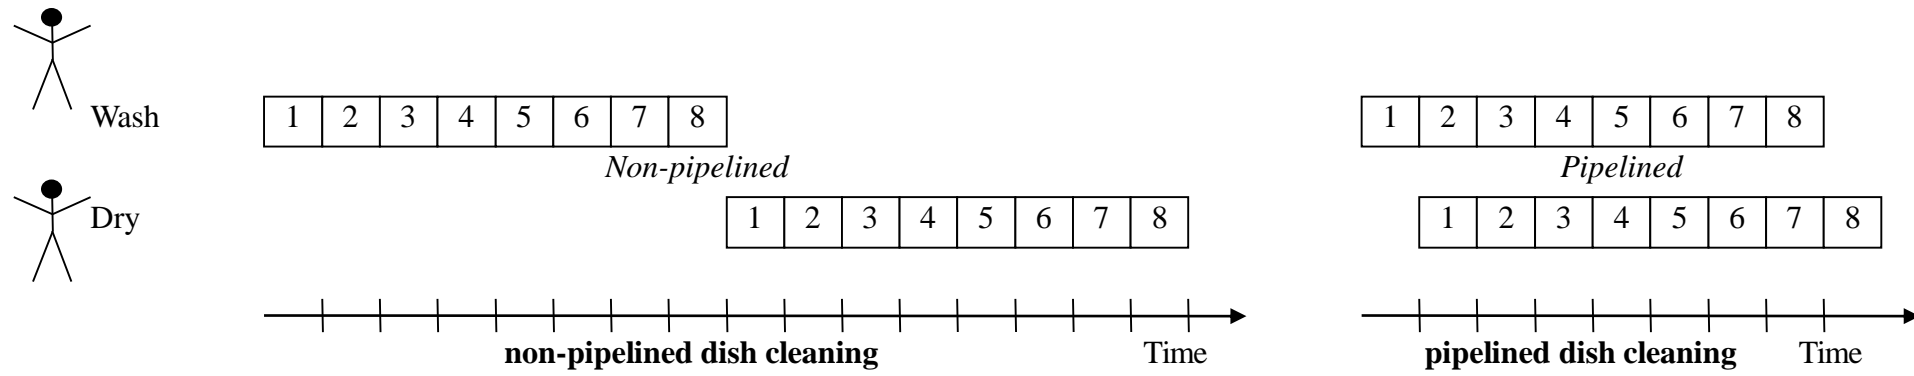


Architectural Considerations

- Clock period
 - Inverse of clock frequency
 - Must be longer than longest register to register delay in entire processor
 - Longest path is called critical path (i.e. from a datapath register through the ALU and back to a datapath register)
 - Memory access is often the longest



Pipelining: Increasing Throughput



Why is branching a problem for pipelining?
What are possible solutions?

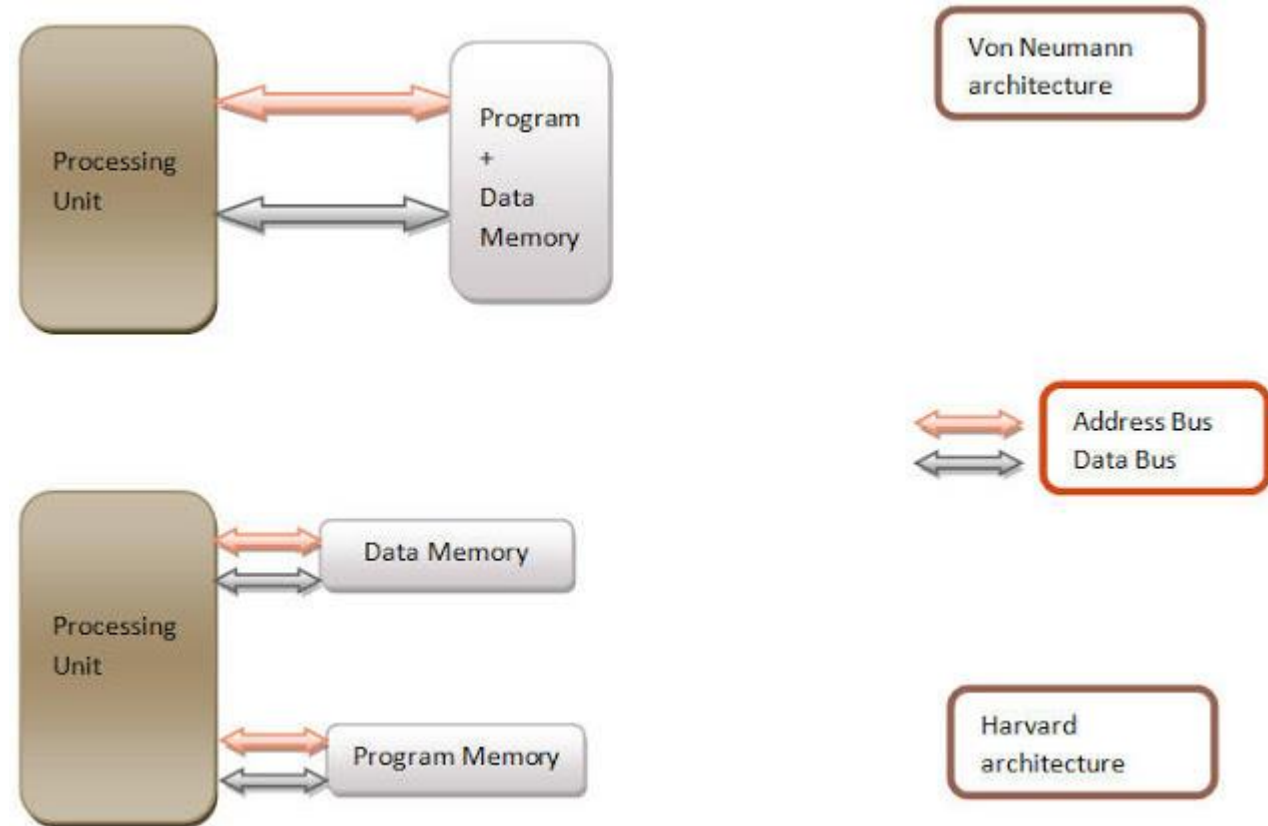
Atmel AVR has 2 stage pipeline
Intel Pentium 4 processors have 20 stage pipelines



Memory Architectures

Memory Architectures

- von Neumann
 - Single data bus which fetches both instructions and data
 - Fewer memory wires, i.e. simpler hardware
 - Bus is bottleneck (von Neumann Bottleneck)
- Harvard
 - Two data buses
 - Different address spaces
 - Simultaneous program and data memory access, i.e. improved performance



AVR uses a Harvard architecture

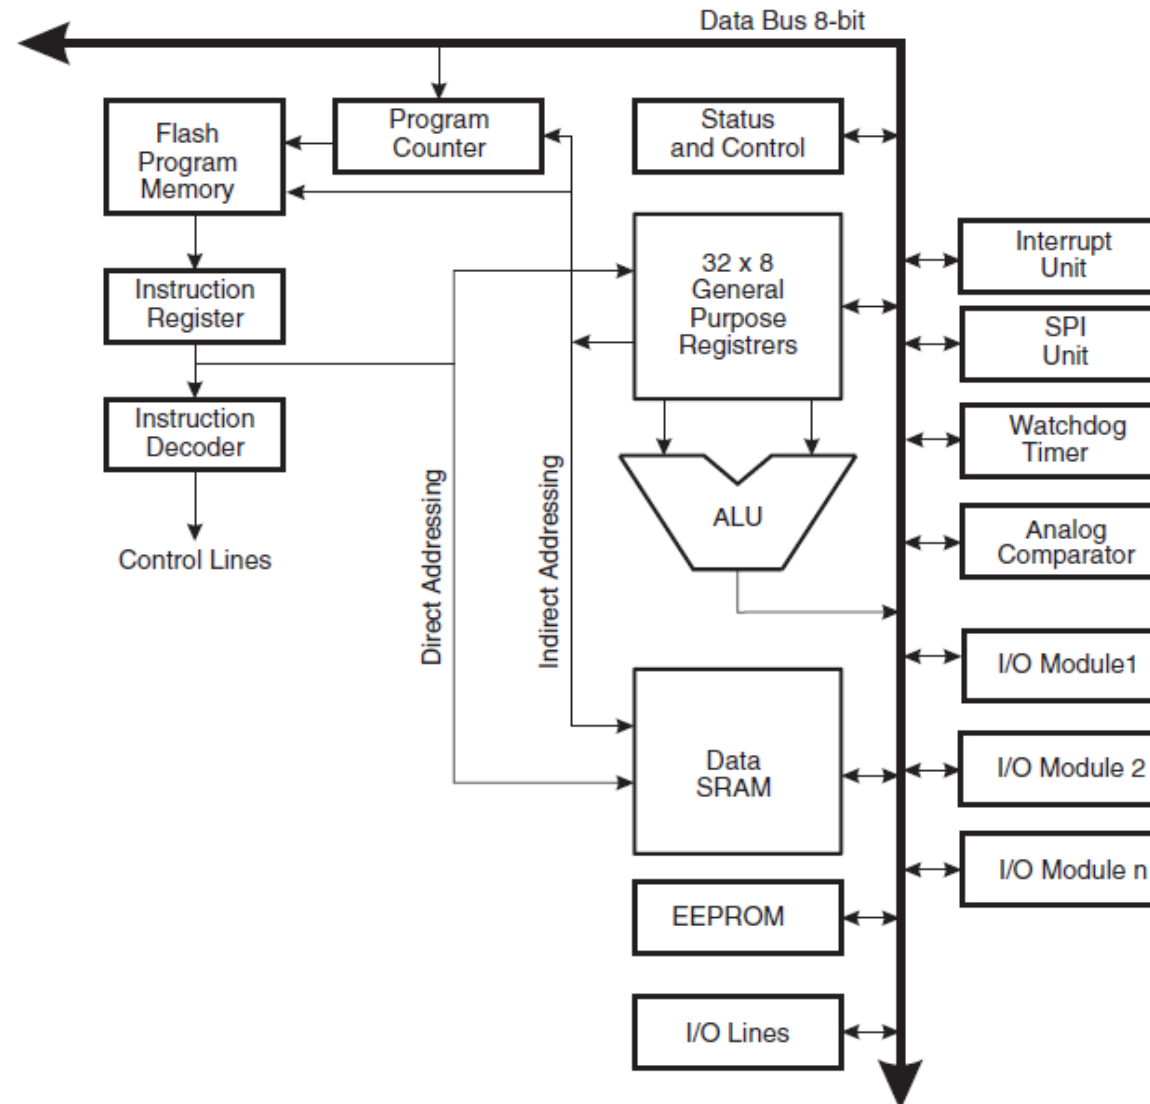
Advantage of Harvard Architecture

- Parallel fetching of data and instructions
- Optimized storage technology for each purpose
 - Program memory as ROM (e.g. Flash)
 - Data memory as RAM
 - Word length can be optimized
 - Data: byte wise access
 - Program: Size of instruction (e.g. 12, 14, 16 Bits)
- Size of address space is doubled
 - Example: 16 Bit addresses: von Neumann $2^{16} = 64$ KBytes, Harvard 128 KBytes

Modified Harvard Architecture

- von Neumann model of stored program computing has advantage over *pure* Harvard machines
 - code can be treated as data, and vice versa
 - reading a program from disk storage and executing it
 - better memory utilization (trade program for data memory)
- Modified Harvard Architecture
 - Variation of Harvard computer architecture that allows contents of instruction memory to be accessed as if it were data, allows code to be generated

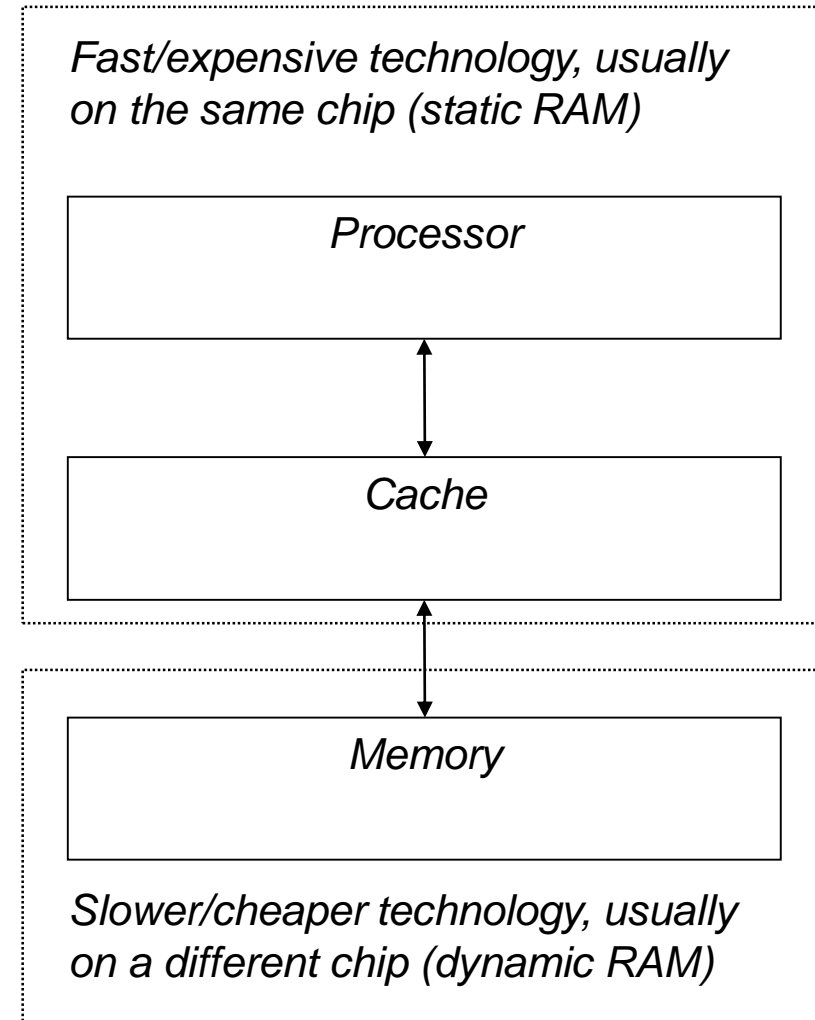
ATmega1281 uses a Harvard Architecture



Source: ATMEL

Cache Memory

- Memory access may be slow
- Cache is small but fast memory close to processor
 - Holds copy of part of memory
 - Hits and misses
 - Caches are used for instructions (I-cache) and data (D-cache)






3

Application Development

Application Development

- Development processor
 - Processor on which we write and debug programs
 - Usually a PC, laptop
- Target processor
 - Processor the program will run on in embedded system
 - Often different from development processor




Toolchain

- Toolchain
 - Set of software development tools that are linked (or chained) together to produce an executable application for a processor
 - For embedded development cross toolchains are used
 - Toolchain runs on a development system of a specific architecture (such as x86) but produce binary code (executables) to run on target architecture (e.g. AVR)
- Optionally, a toolchain may contain other tools such as a debugger
- Various open source projects that comprise entire toolchain are available

Tools

- Preprocessor
 - Processes source code (inclusion of header files, macro expansions, conditional compilation, etc.)
- Compiler
 - Translates human readable code into assembly language for a particular processor
- Assembler
 - Translates assembly language into opcodes. Produces an object file
- Linker
 - Organizes object files, necessary libraries, and other data and produces a relocatable file
- Locator
 - Takes relocatable file and information about system's memory and produces an executable, e.g. transform relocatable addresses into absolute addresses
- gcc (GNU Compiler Collection) takes care of *all steps* at once!

Embedded Software Development Process



Tools: Embedded System Specifics

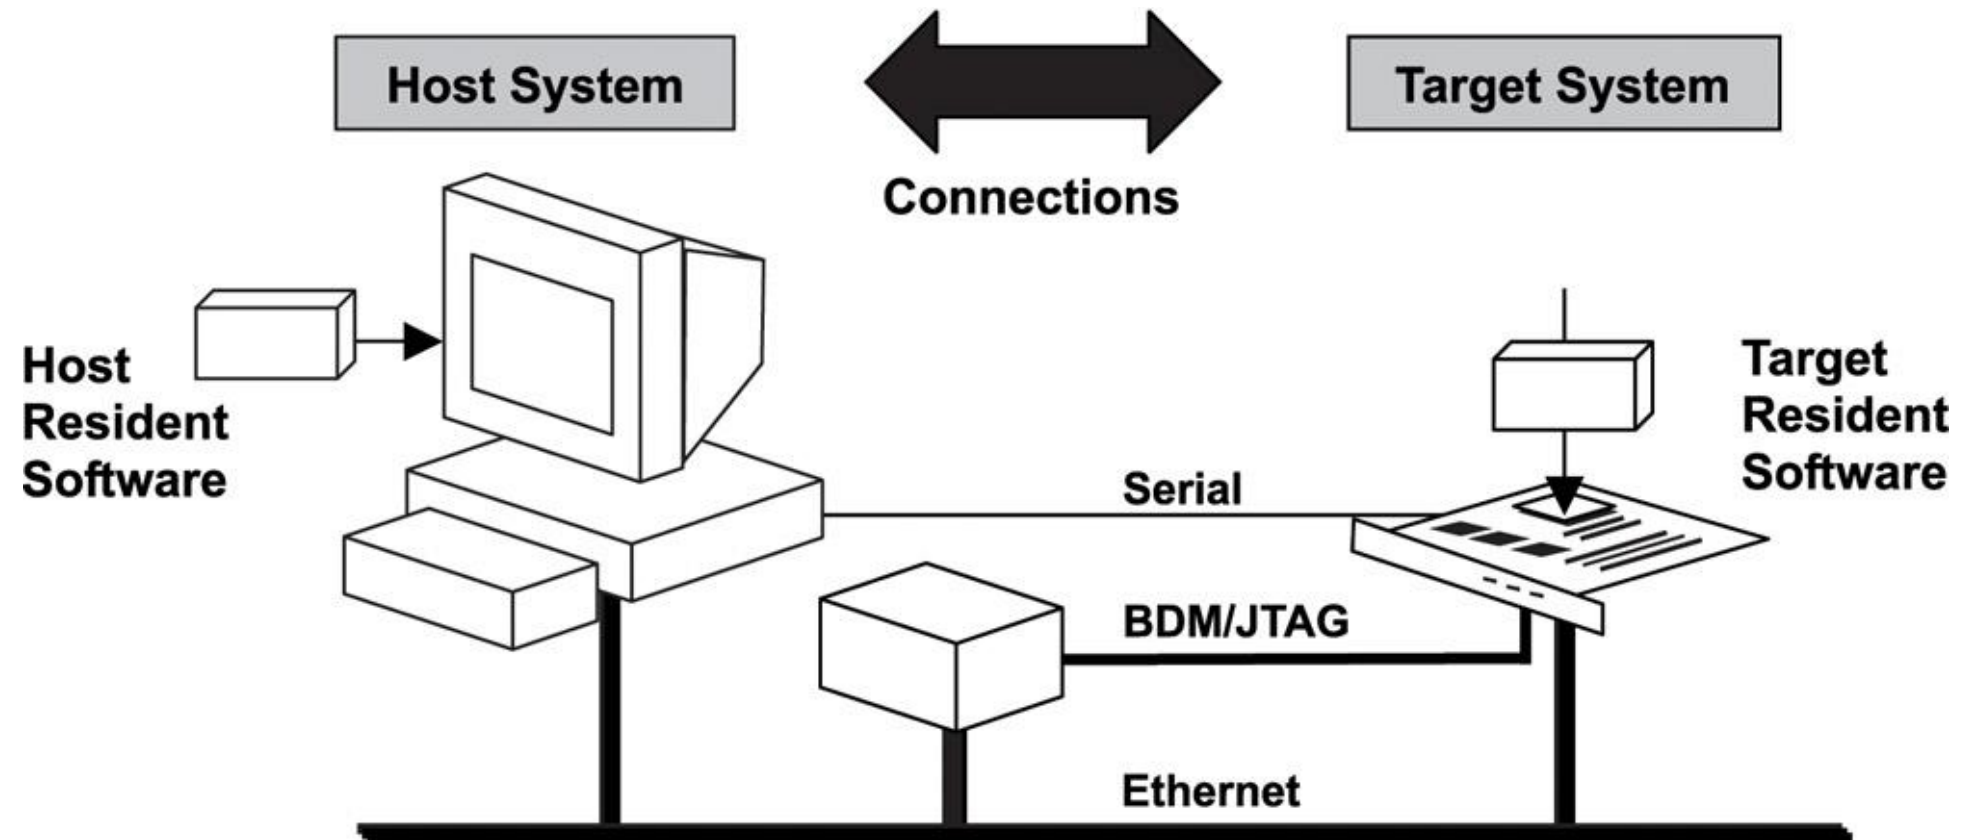
- Tools run on development computer, not on embedded computer
- Compiler
 - Has to know specific hardware, optimizes for size & speed
- Assembler
 - Produces “startup code”; not inserted automatically as in general purpose computers (i.e., programmer needs to compile it independently)
- Linker
 - Needs correct libraries
- Locator
 - Needs programmer input for information about memory
- Complete embedded tool chain often integrated into Integrated Development Environment (IDE)
 - e.g. Eclipse, AVR Studio

Moving program onto embedded system

- Compiled program (image) needs to get onto embedded system
- Methods to load image:
 - In-System Programming (ISP)
 - Programmed while installed in a complete system, rather than requiring chip to be programmed prior to installing it into system
 - Entire image installed into EEPROM or flash memory from serial port of PC (e.g. using SPI)
 - Downloading image through a JTAG or BDM interface
 - Bootloader
 - Data transfer utility program on host, as well as a target loader on embedded system
 - Loads programs onto system over a serial line (e.g. RS-232)
 - Requires no special hardware



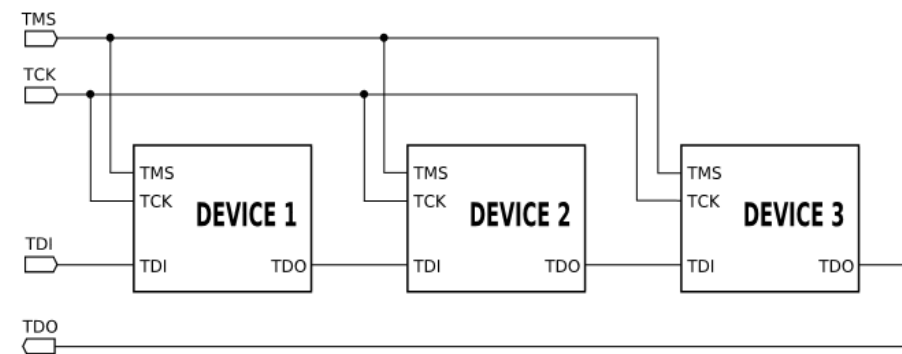
Cross-platform development environment



Quelle: book.opensourceproject.org.cn

Joint Test Action Group (JTAG)

- JTAG: Originally interface used for testing PCBs after assembly
- Chip's IO lines can be controlled and read via JTAG ports allowing a board test sequence to be performed
- JTAG also allows device programmer hardware to transfer data into internal non-volatile memory
- Advantage of JTAG
 - Devices can be daisy chained: One JTAG port can serve to program/debug/test multiple devices





Programming

Programming

- Programmer doesn't need detailed understanding of architecture
 - Instead, needs to know what instructions can be executed
- Three levels of instructions:
 1. Machine language
 2. Assembly level
 3. High-level structured languages (C, C++, Ada, Java, etc.)
- Today mainly high-level languages are used
 - They do not depend on hardware or abilities of CPU
 - But, some assembly level programming may still be necessary
- High-level languages provide device driver APIs
 - Driver: program communicating with and controls (*drives*) a device

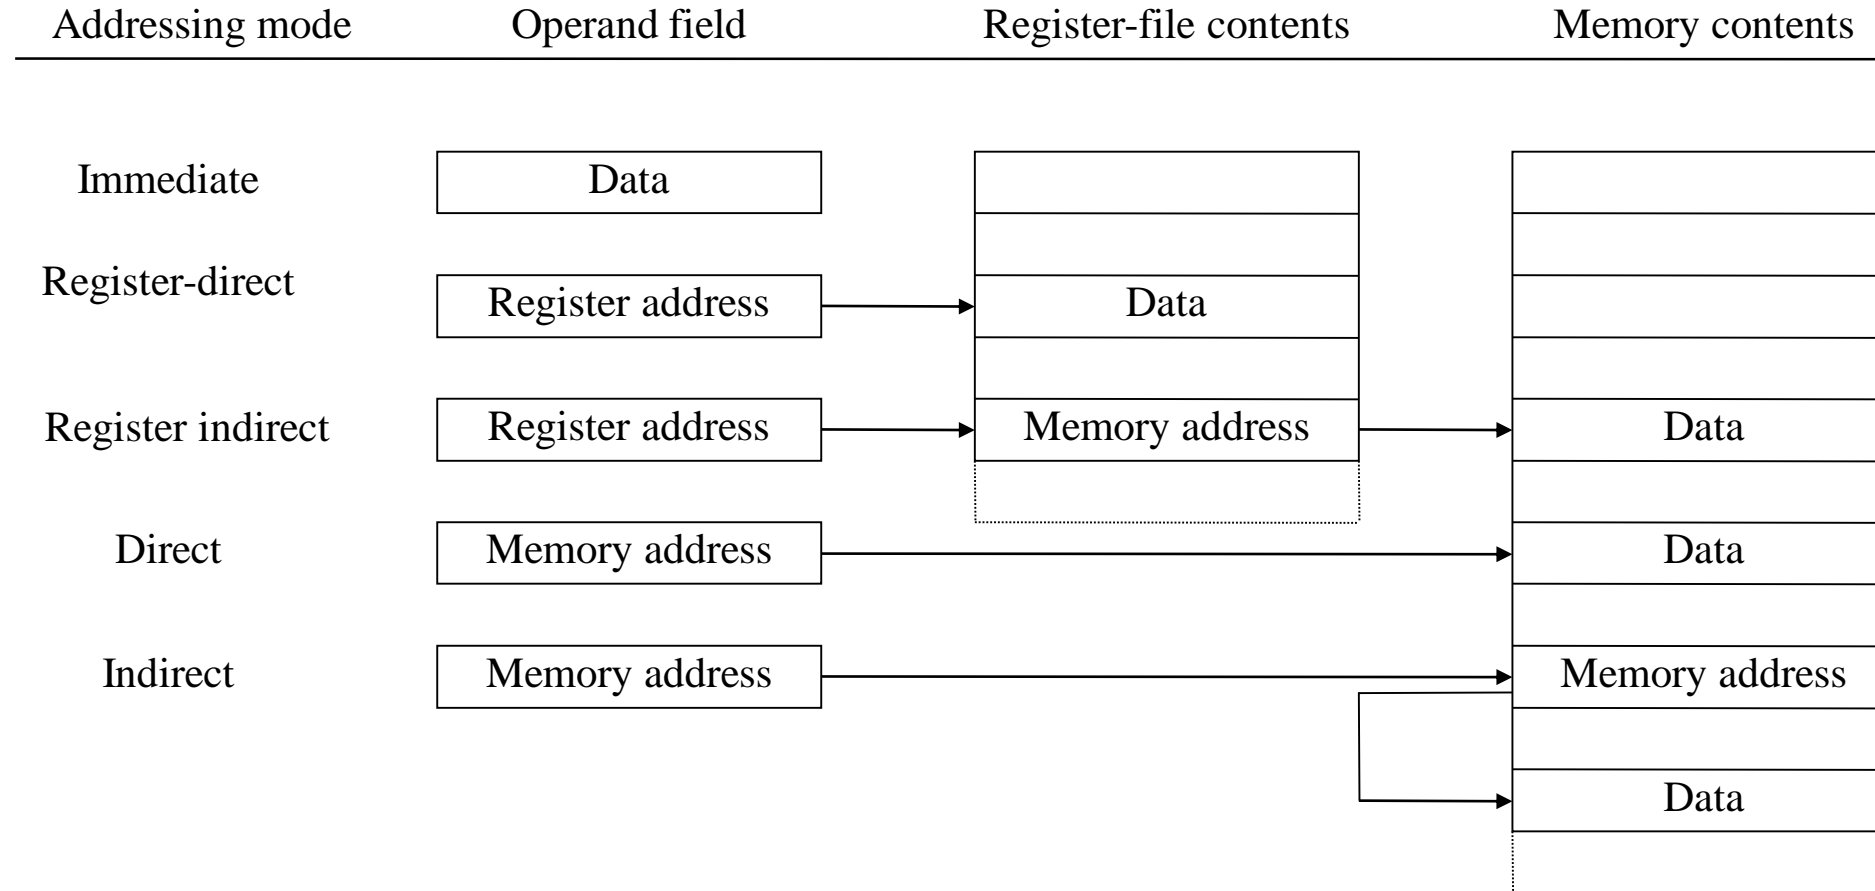
Assembly-Level Instructions

Instruction 1	opcode	operand1	operand2
Instruction 2	opcode	operand1	operand2
Instruction 3	opcode	operand1	operand2
Instruction 4	opcode	operand1	operand2
...			

- Structure
 - opcode field
 - operand fields
 - number varies (source and destination)
 - addressing modes

- Instruction Set: Legal set of instructions of processor
 - Data transfer: memory/register, register/register, I/O, etc.
 - Arithmetic/logical: move register through ALU
 - Branches: determine next PC value when not just PC+1
- AVR instructions are 16 or 32 bits wide

Addressing Modes



A Simple Instruction Set

Assembly instruct.	First byte		Second byte	Operation
MOV Rn, direct	0000	Rn	direct	$Rn = M(\text{direct})$
MOV direct, Rn	0001	Rn	direct	$M(\text{direct}) = Rn$
MOV @Rn, Rm	0010	Rn	Rm	$M(Rn) = Rm$
MOV Rn, #immed.	0011	Rn	immediate	$Rn = \text{immediate}$
ADD Rn, Rm	0100	Rn	Rm	$Rn = Rn + Rm$
SUB Rn, Rm	0101	Rn	Rm	$Rn = Rn - Rm$
JZ Rn, relative	0110	Rn	relative	$PC = PC + \text{relative}$ (only if Rn is 0)
	opcode		operands	

Destination is (almost) always given first, then source!

Sample Programs

C program

```
int total = 0;
for (int i=10; i!=0; i--)
    total += i;
// next instructions...
```

Equivalent assembly program

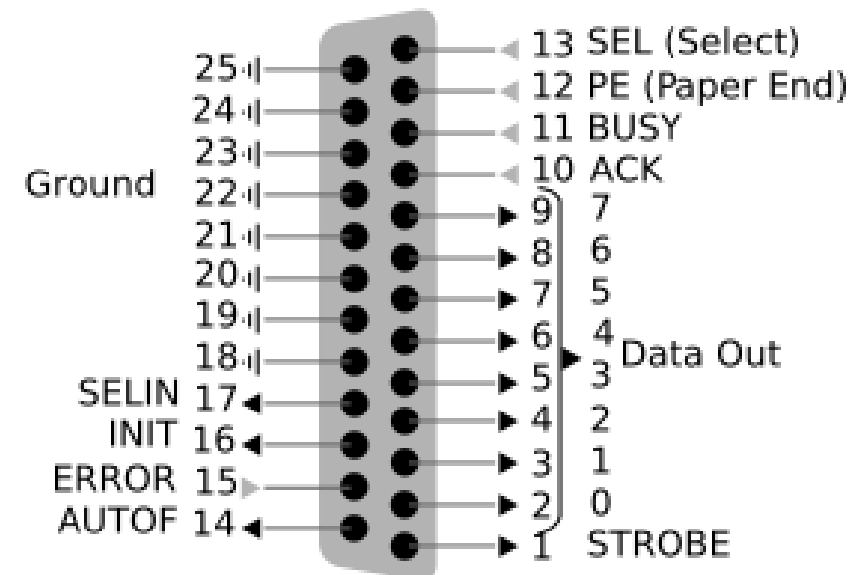
```
0      MOV R0, #0;        // total = 0
1      MOV R1, #10;       // i = 10
2      MOV R2, #1;        // constant 1
3      MOV R3, #0;        // constant 0
Loop:  JZ R1, Next;        // Done if i=0
5      ADD R0, R1;         // total += i
6      SUB R1, R2;         // i--
7      JZ R3, Loop;        // Jump always
Next:  // next instructions...
```

■ Try some others

- Handshake: Wait until the value of M[254] is not 0, set M[255] to 1, wait until M[254] is 0, set M[255] to 0 (assume those locations are ports)
- (Harder) Count the occurrences of zeros in an array stored in memory locations 100 through 199


LPT (Line Print Terminal)

- An LPT port (parallel port) has an
 - 8-bit data bus
 - four pins for control output (Strobe, Linefeed, Initialize, and Select In), and
 - five more for control input (ACK, Busy, Select, Error, and Paper End)
- Data transfer speed
 - 12,000 kbit/s



Example: Parallel Port Driver

LPT Connection Pin	I/O Direction	Register Address
1	Output	0 th bit of register #2
2-9	Output	0 th - 7 th bit of register #0
10,11,12,13,15	Input	6,7,5,4,3 th bit of register #1
14,16,17	Output	1,2,3 th bit of register #2



- Using assembly language programming we can configure a PC parallel port to perform digital I/O
 - write and read to three special registers (#0, #2 for output, #1 for input) to accomplish this
 - table provides list of parallel port connector pins and corresponding register location
- Example: parallel port monitors the input switch and turns the LED on/off accordingly

Parallel Port Example

```
; This x86 program consists of a sub-routine reading the state
; of the input pin, determining the on/off state of our switch
; and asserts the output pin, turning LED on/off accordingly
```

```
.386
```

```
checkPort  proc
```

```
push  al           ; save the content
push  dx           ; save the content
mov   dx, 0x3BC + 1 ; base + 1 for register #1
in    al, dx       ; read register #1
and   al, 0x10     ; mask out all but bit # 4
cmp   al, 0        ; is it 0?
jne   SwitchOn     ; if not, turn LED on
```

```
SwitchOff:
```

```
mov   dx, 0x3BC + 0 ; base + 0 for register #0
in    al, dx        ; read current state of port
and   al, 0xfe      ; clear first bit (masking)
out   dx, al        ; write it out to the port
jmp   Done          ; we are done
```

```
SwitchOn:
```


```
mov   dx, 0x3BC + 0 ; base + 0 for register #0
in    al, dx        ; read current state of port
or    al, 01h       ; set first bit (masking)
out   dx, al        ; write it out to port
```

```
Done:
```

```
pop   dx           ; restore the content
pop   al           ; restore the content
checkPort  endp
```

```
extern "C" checkPort(void); // defined in assembly
```

```
void main(void) {
    while (1) {
        checkPort();
    }
}
```




LPT Connection Pin	I/O Direction	Register Address
1	Output	0 th bit of register #2
2-9	Output	0 th – 7 th bit of register #0
10,11,12,13,15	Input	6,7,5,4,3 th bit of register #1
14,16,17	Output	1,2,3 th bit of register #2

Address of register #0 is 0x3BC

Parallel Port Example

- General register al, dx are saved/restored
- Different notation for hex values: 72h, 72H, 0x72, \$72, 72₁₆;
- Switch 0 means switch off
- In/out instructions read/write internal registers
 - Operands: address and data
 - In: content of 8-bit operand will be written to addressed register
 - Out: content of addressed 8-bit register will be read into operand (operation direction right to left)
- Address is calculated by adding address of device (0x3BC) to address of register
- `camp al, imm8` : Compares first operand with second operand and sets status flags in EFLAGS register according to result
- `cmp` instruction typically used in conjunction with conditional jump

Implementation of an Embedded Program



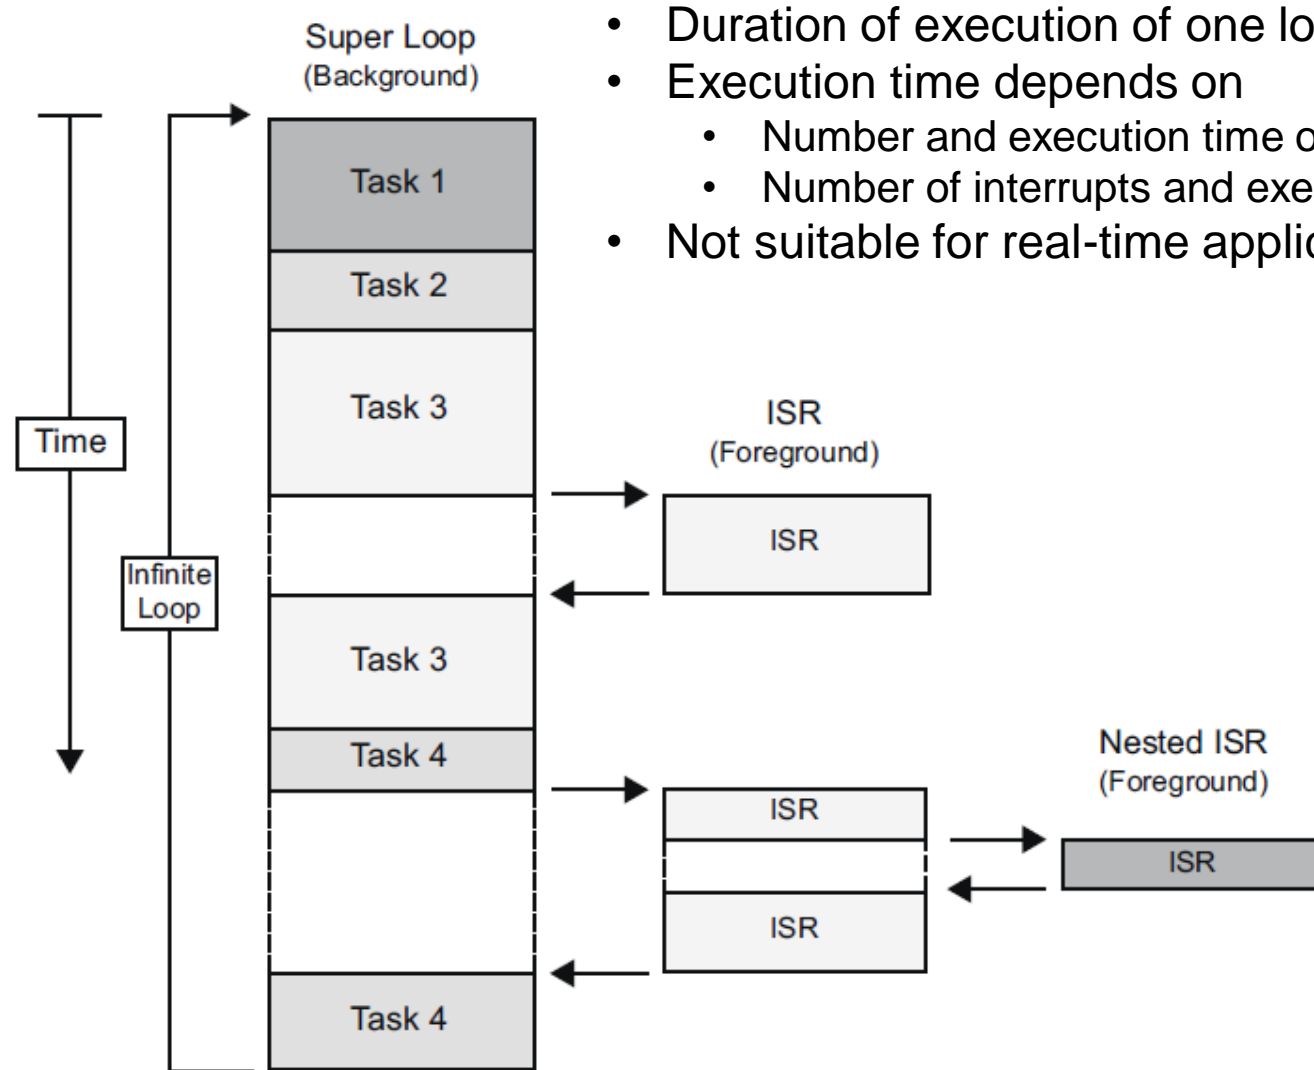
- Super Loop:
 - runs forever
 - is used for tasks of low/medium complexity
 - cannot be used to periodically execute a function at fixed intervals
 - only interrupts intercept the super loop

Super Loop

```
int main(void) {  
    initialize();           // executed once  
  
    // super loop starts here  
    // tasks are executed repetitively  
    while (1) {  
        task1();           // e.g. read temperature  
        task2();           // e.g. process temperature  
        task3();           // e.g. display temperature  
        delay_loop();      // to slow down the process  
    }  
}
```

- In delay processor may transit into low power mode
- Alternative: Scheduler, operating system (OS)

Super Loop with Interrupts



- Duration of execution of one loop is unknown
- Execution time depends on
 - Number and execution time of tasks
 - Number of interrupts and execution time of ISRs
- Not suitable for real-time applications

Operating System

- Software layer providing low-level services to a program
 - Loading & executing programs, process
 - Sharing, allocate and manage system resources
 - File management, disk access, keyboard/display interfacing
 - Scheduling multiple programs for execution
 - Program makes system calls to OS
 - System Call:
Mechanism for application to invoke the OS via software interrupt

```
DB file_name "out.txt"      -- store file name

MOV R0, 1324                -- system call "open" id
MOV R1, file_name           -- address of file-name
INT 34                      -- cause a system call
JZ  R0, L1                  -- if zero -> error

    ... read the file
JMP L2                      -- bypass error cond.
L1:
    ... handle the error

L2:
    ...
```




5

Running a Program

Running a Program

- If development processor is different than target, how can we run our compiled code? Two options:
 - Download to target processor
 - Simulate
- Simulation
 - One method: Hardware description language
 - But slow, not always available
 - Another method: *Instruction set simulator (ISS)*
 - Runs on development processor, but executes instructions of target processor

A Simple Instruction Set

Assembly instruct.	First byte		Second byte	Operation
MOV Rn, direct	0000	Rn	direct	$Rn = M(\text{direct})$
MOV direct, Rn	0001	Rn	direct	$M(\text{direct}) = Rn$
MOV @Rn, Rm	0010	Rn	Rm	$M(Rn) = Rm$
MOV Rn, #immed.	0011	Rn	immediate	$Rn = \text{immediate}$
ADD Rn, Rm	0100	Rn	Rm	$Rn = Rn + Rm$
SUB Rn, Rm	0101	Rn	Rm	$Rn = Rn - Rm$
JZ Rn, relative	0110	Rn	relative	$PC = PC + \text{relative}$ (only if Rn is 0)
MOV Rn, @Rm	0111	Rn	Rm	$Rn = M(Rm)$
	opcode		operands	

Destination is always given first, then source!

Example for Instruction Set Simulator

```
#include <stdio.h>
typedef struct {
    uint8_t first_byte, second_byte;
} instruction_t;

instruction_t program[PROGRAMSIZE]; // instruction memory
uint8_t memory[MEMORYSIZE];        // data memory
uint8_t reg[NUM_REGISTERS];         // register file

int8_t run_program(int16_t num_instr) {
    int16_t pc = -1;
    uint8_t fb, sb, fb_low, sb_high;

    while( ++pc < num_instr ) {
        fb = program[pc].first_byte;
        sb = program[pc].second_byte;
        fb_low = fb & 0x0f; // Rn
        sb_high = sb >> 4;  // Rm
        switch( fb >> 4 ) { // opcode
            case 0: reg[fb_low] = memory[sb]; break;
            case 1: memory[sb] = reg[fb_low]; break;
            case 2: memory[reg[fb_low]] =
                    reg[sb_high]; break;
            case 3: reg[fb_low] = sb; break;
            case 4: reg[fb_low] += reg[sb_high]; break;
            case 5: reg[fb_low] -= reg[sb_high]; break;
            case 6: if (reg[fb_low] == 0) pc += sb; break;
```

```
            case 7: reg[fb_low] = memory[reg[sb_high]]; break
            default: return -1;
        }
    }
    return 0;
}

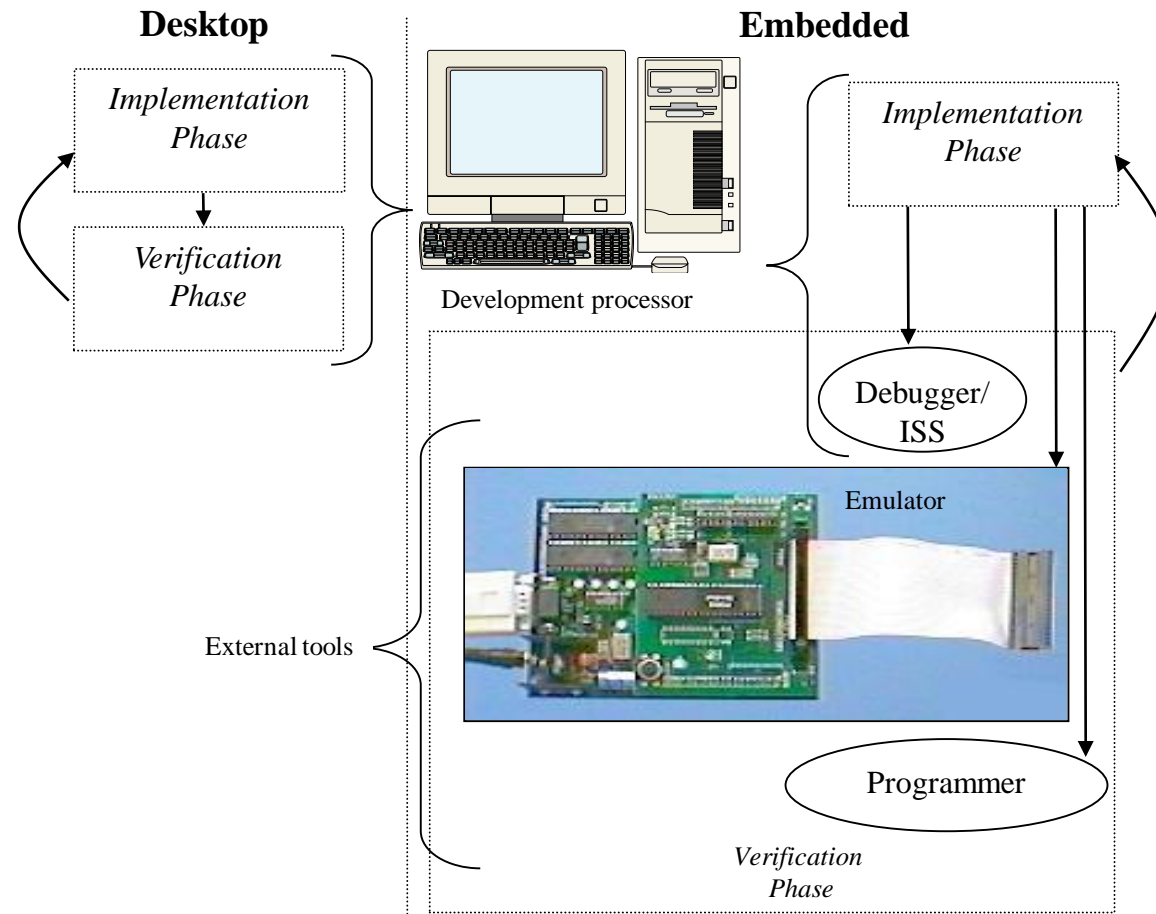
void print_memory_contents() {
    // Output of content of memory
}

int16_t main(int argc, char *argv[]) {
    FILE * ifs;

    if (argc != 2 || (ifs = fopen(argv[1], "rb") == NULL ) {
        return -1;
    }

    if (run_program(fread(program, sizeof(instruction),
        sizeof(program)/sizeof(instruction), ifs)) == 0)
    {
        print_memory_contents();
        return 0;
    } else {
        return -1;
    }
}
```

Testing and Debugging



- **ISS**
 - Gives control over time – set breakpoints, look at register values, set values, step-by-step execution, ...
 - But, doesn't interact with real environment
- **Download to board**
 - Use device programmer
 - Runs in real environment, but not controllable
- **Compromise: emulator**
 - Runs in real environment, at speed or near
 - Supports some controllability from the PC

Chapter Summary

- General-purpose processors
 - Good performance, low NRE, flexible
- Controller, datapath, and memory
- Structured languages prevail
 - But some assembly level programming still necessary
- Many tools available
 - Including instruction-set simulators, in-circuit emulators

SES

Chapter 2: General-Purpose Processors

Prof. Dr.-Ing. Bernd-Christian Renner

