

# SES

## Chapter 6: Operating Systems for Embedded Systems

Prof. Dr.-Ing. Bernd-Christian Renner



# Contents

1. Schedulers
  - a) Round Robin Scheduler
  - b) Periodic Scheduler
  - c) Cooperative Scheduler
  - d) Preemptive Scheduler
2. Multitasking Operating Systems
3. Concurrent Task Model
4. Communication Among Tasks



# Implementing Embedded Systems

- Embedded systems are often reactive systems
  - they respond to external events
  - often require real-time response
- **Simplest structure for reactive programs: Super-loop**
  - Functional related operations are grouped into tasks which are strung together in a large endlessly executing loop
  - Each task
    - is coded as a separate block of code with a single start and a single end point (e.g. as a single function)
    - first checks for a relevant state change and if necessary executes code
- **Completely deterministic model**, i.e. response time can be calculated (if execution times are fixed)



# Schedulers

# Overview

- Higher level:
  - Scheduler allows tasks to be called (periodically or sporadically)
- Lower level:
  - Scheduler can be viewed as a single timer interrupt service routine that is shared between many different tasks
    - i.e. only one timer needs to be initialized, any changes to the timing generally requires only one function to be altered
- Scheduler is independent of nature/number of tasks
  - Advantage: Scheduler has to be written only once
- Scheduler: Beginning of operating system

# Round Robin Scheduler

- Round Robin Scheduler (RRS)
  - Scheduler maintains a list of all active tasks (task list)
  - Tasks operate on shared data
  - RRS takes care of switching from one task to the next
  - Programmer has to initially tell scheduler which tasks are to run, but after that scheduler takes control
  - All tasks run to completion (not preemptive)
- Design of a simple RRS
  - Implement scheduler
  - Declare prototypes for tasks
  - Implement tasks

# Tasks

- Each task must have the same signature
- Generic case
  - Return type `void`
  - One parameter of type `void *` (can also be used to return values)

- Tasks are stored in array

```
void (*queue[TASKCOUNT])(void *);
```

- Example: Three tasks (parameter is address of shared data)

```
void get(void * aNumber);           // input task
void increment(void * aNumber);     // computation task
void display(void * aNumber);       // output task
```

# Implementation of Tasks

```
void get(void * aNumber) {                                //input task
    int * intPtr = (int *)aNumber;
    printf("Enter a number: 0..9");
    *intPtr = getchar();
    getchar();
    *intPtr -= '0';
    return;
}
```

```
void increment(void * aNumber) {                            // computation task
    int * intPtr = (int *)aNumber;
    (*intPtr)++;
    return;
}
```

```
void display(void * aNumber) {                               // output task
    printf("The result is: %d\n", *(int *)aNumber);
    return;
}
```




# RRS: Code for Scheduler

```
void main() {  
    int i = 0;                // queue index  
    int data;                 // shared data  
    int * intPtr = &data;    // pointer to data  
  
    void (*queue[TASKCOUNT])(void *); // task queue  
  
    queue[0] = get;  
    queue[1] = increment;  
    queue[2] = display;  
  
    while (1) {                // scheduler  
        queue[i]((void*) intPtr); // dispatch task  
        i = (i+1) % TASKCOUNT;  
    }  
    return;  
}
```

# Example: Central-heating controller

- 3 Tasks
- Fixed order
- Round Robin Scheduler is simple solution



# Limitations

- Tasks are run in a fixed order
  - Tasks cannot be skipped
  - No priorities
- No control over timing of tasks
  - A task has to wait until all preceding tasks have finished
  - Execution period of a task cannot be predicted
    - task execution times may vary
- Operates at 'full power' at all times
- No ad-hoc execution of a task
- Extensions:
  - Dynamically add, remove, stop, and resume tasks
  - Execute tasks upon events
  - Execute tasks at specific times

# Alternative: Event Driven Loop

- Tasks are executed upon events
- All tasks are implemented as ISRs
- Main program:
  - infinite loop interrupted when event occurs
  - flow of control jumps to associated ISR, where designated task is executed
  - afterwards flow resumes infinite loop
- More difficult to analyze due to asynchronous events
- Often used with priorities to make it more deterministic

```
global variable declarations

ISR set up
function prototypes
void main (void)
{
    local variable declarations
    while(1);           // task loop
}
ISRs
function definitions
```



# Periodic Scheduler

- Tasks are executed periodically
- Example:
  - Tasks are finite state machines (FSM), i.e., each task is implemented by a function with signature  
`uint8_t (*func) (uint8_t)`
  - Each task can have its own execution period (multi-rate)

- Data structure for tasks

```
typedef struct {  
    uint8_t state;  
    uint32_t period;  
    uint32_t elapsedTime; // time since last execution  
    uint8_t (*func) (uint8_t)  
} task_t;
```

# Periodic Scheduler

- Initialization

```
task_t task_1, task_2, ...;
task_t * tasks[] = {&task1, &task2, ...}
const uint8_t numTasks = sizeof(tasks)/sizeof(tasks[0]);

//Init each task
task_1.state = -1; // initial state of FSM
task_1.period = ...;
task_1.elapsedTime = task_1.period;
// all tasks are immediately executed
task_1.func = func1;
...

//ISR for timer
volatile uint8_t timerFlag = 0;

void timerISR() {
    timerFlag = 1;
}
```

# Periodic Scheduler

- Code for Scheduler

```
timerSet(basePeriod);  
timerOn();
```

```
while (1) {  
    for (i=0; i < numTasks; i++) {  
        if (tasks[i]->elapsedTime == tasks[i]->period) {  
            task[i]->state =  
                tasks[i]->func(tasks[i]->state);  
            tasks[i]->elapsedTime = 0;  
        }  
        tasks[i]->elapsedTime += basePeriod;  
    }  
    while (timerFlag == 0);  
    timerFlag = 0;  
}
```

All tasks together must  
complete within basePeriod


# Aperiodic tasks

- Not all tasks are periodic tasks
  - Example: A system has to monitor an input source for occurrence of a special value
- Possible solution: Periodic polling
  - Con: High microcontroller usage
- Alternative:
  - Some microcontrollers come with special hardware to detect changes on an input and call an ISR
  - Example: UART provides such events
- New concept: Asynchronous finite state machine
  - A task is triggered by an event and not executed periodically



# Scheduler for periodic and aperiodic tasks

- Extend data structure for task by  
`volatile uint8_t active;`
- Model each monitoring task as



- For each monitoring task provide separate ISR

```

void eventA0_ISR() {
    tasks[i]->active = 1;
    tasks[i]->elapsedTime = tasks[i]->period;
    eventFlag = 1;
}
  
```

- Task `tasks[i]` handles event `A0`

# Main Loop

```
volatile unit8_t eventFlag = 0;
while (1) {
    for (i = 0; i < numTasks; i++) {
        if (tasks[i]->active) {
            if (tasks[i]->elapsedTime == tasks[i]->period) {
                task[i]->state = tasks[i]->func(tasks[i]->state);
                tasks[i]->elapsedTime = 0;
                if (tasks[i] is non-periodic task) {
                    tasks[i]->active = 0;
                }
            }
            tasks[i]->elapsedTime += basePeriod;
        }
    }
    while (timerFlag == 0 && eventFlag == 0) { /* busy wait */ };
    timerFlag = 0;
    eventFlag = 0;
}
```

# Advanced Schedulers

- Schedulers considered so far let each task run to completion before next task is started
- Alternative:
  - Execution of a task can be halted and another task begins execution
  - At a later point execution of first task is resumed
- Realizations:
  1. Cooperative scheduler
    - **Tasks** decide themselves when to halt and yield to another task
  2. Preemptive scheduler
    - **Scheduler** decides when to halt current task and resume another task
- Changing of running task is called *context switch*

# Advanced Schedulers

- Dispatcher:
  - Component that is responsible for remembering yielding task's whereabouts (called task's execution context) and starting up next task from task set
  - Execution context is important for resuming the task
    - contents of all registers (including program counter) and stack
  - Tasks of dispatcher
    - Save execution context of current task in RAM
    - Load execution context of next task from RAM into registers and restore stack
- Scheduler:
  - Decides which task to run next
  - May use priorities
  - Requires data structure to maintain tasks



# Cooperative Schedulers

- Operation of a cooperative scheduler:
  - Task executes a special instruction that lets it yield control back to scheduler (implemented as a software interrupt)
  - Scheduler selects task to run next
  - Scheduler advises dispatcher to do a context switch
    - Next task starts off at exactly where it left off
- When does a task yield?
  - it has done all it can for now
  - taken up enough processor time
- Problems with cooperative schedulers
  - What happens if task *forgets* to yield?
  - What happens if a task crashes?

# Preemptive Schedulers

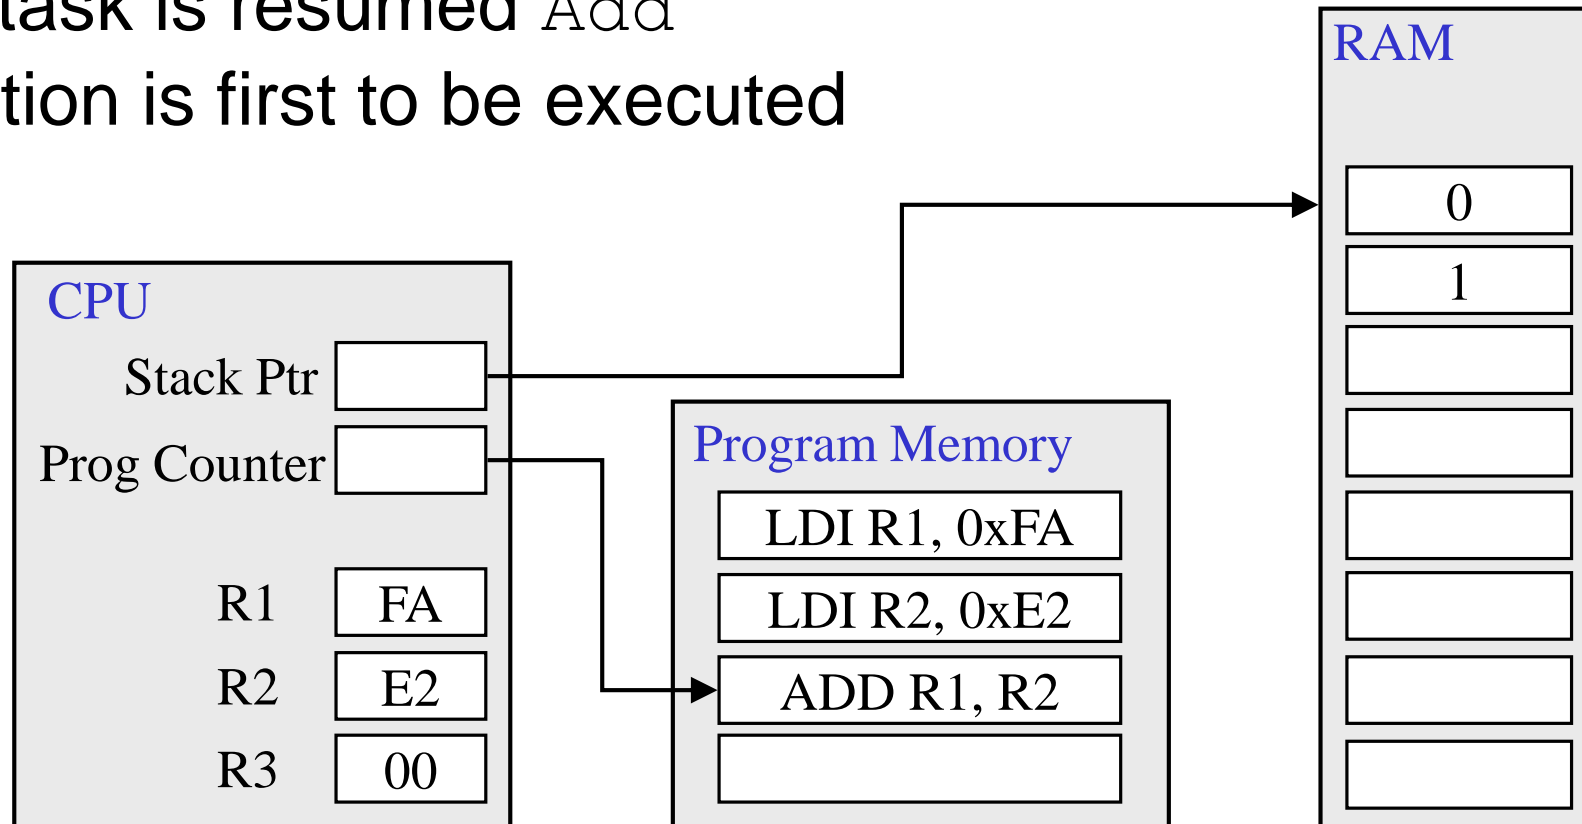
- In preemptive scheduling tasks can be **forced** to interrupt, i.e. tasks are not yielding voluntarily
- Context switches are triggered by events
  - externally generated by hardware inputs
    - e.g. reading from a UART
  - internally generated by hardware timers
- Interruption of tasks is transparent for users
- Preemptive schedulers also require a dispatcher, same duties as with cooperative scheduler

# Execution Context

- Task does **not** know when it is going to be suspended or resumed
- While task is suspended other tasks will execute and may modify register values
- It is essential that upon resumption a task has a context identical to that immediately prior to its suspension
- Context switch is based on periodic timer
  - Timer ISR checks if a context switch must be performed
  - Context switch is implemented by dispatcher
- Dispatcher realized as ISR

# Example

- Task being suspended is immediately before executing an instruction adding values contained in registers R1 and R2. When task is resumed `Add` instruction is first to be executed





# AVR Context

- AVR Context:
  - 32 general purpose registers
  - Status register
    - Its value affects instruction execution, and must be preserved across context switches
  - Program counter
    - Upon resumption, a task must continue execution from the instruction that was about to be executed immediately prior to its suspension
  - Two stack pointer registers
  - Stack

# AVR Execution Context

## General Purpose Registers



## Status




## Programm Counter



## Stack Pointer



## Stack



# ISR and Execution Context

- Code of ISR must save and restore context
- Compiler generated ISR ensures that **every register that gets modified by ISR** is restored to its original value when ISR exits
  - Upon ISR completion:  
Always return to code when interrupt occurred
  - Context switch is different:  
Compiler cannot make assumptions as to when context switch will happen, and therefore cannot optimize which registers require saving and which don't (any code could be executed next)
- Consequence:  
Entire context must be saved and context of next task must be loaded

# Saving the Context

- Generated ISR code forces a *return from interrupt* instruction (RETI) to be used in place of the 'return' instruction (RET) that would otherwise be used
  - AVR microcontroller disables interrupts upon entering an ISR and RETI instruction is required to re-enable them on exiting
- Context switch requires Dispatcher to explicitly save all registers on entering ISR, but doing so would result in some registers being saved twice
  - once by compiler generated code and again by dispatcher code
- Solution: `naked` attribute
  - prevents compiler from generating any function entry or exit code so this must be added explicitly
  - gives application code complete control over when and how context is saved/restored

# Saving the Context

- Context is pushed on stack
  - program counter is pushed by ISR automatically
  - r0 to r31, SREG, and stack pointer must be pushed manually
- Each task requires its own stack, when a task swap is performed, pointers to corresponding stack are restored
- Allocate for each task to be executed (how many?) memory (how much?) to host task's stack and to store stack pointer
- This structure is called **task control block** (TCB)
- When task is
  - preempted: push context on stack, write stack pointer variable
  - resumed: restore stack pointer, pop and restore context from stack
- Memory for each stack is fixed, so beware of stack overflow!

# Context Switching ISR

```
void IST_Tick(void) __attribute__((signal, naked));  
void IST_Tick(void)  
{  
    /*  
        Save execution context  
        Implement clock tick  
        Restore new execution context  
    */  
    asm volatile ("reti");  
}
```

## AVR Solution:

```
ISR(TIMER1_COMPA_vect, ISR_NAKED) {
```



# Saving the Context

```
asm volatile (  
    "push r0 \n\t"  
    "in r0, __SREG__ \n\t" //status register moved into R0 to be saved  
    "cli \n\t" //dis. inter. (when called outside ISR, cop. sched.)  
    "push r0 \n\t"  
    "push r1 \n\t"  
    "clr r1 \n\t"          //compiler assumes that r1 is 0  
    "push r2 \n\t"  
    "push r3 \n\t"  
    "...."  
    "push r31 \n\t"       // saving task execution context finished  
    "lds r26, currentTCB \n\t" // load address to which SP is saved  
    "lds r27, currentTCB + 1 \n\t"  
    "in r0, __SP_L__ \n\t" // save stack pointer  
    "st x+, r0 \n\t"  
    "in r0, __SP_H__ \n\t"  
    "st x, r0 \n\t"  
);
```

- program counter automatically pushed on stack
- currentTCB, currentTCB + 1: address from where task's stack pointer can be retrieved
- st x+,r0 stores r0 into location pointed to by X register (16 bits), increments X register

# Restoring the Context

```
asm volatile (  
    "lds r26, currentTCB \n\t" // currentTCB holds address of SP  
    "lds r27, currentTCB + 1 \n\t"  
    "ld r28, x+ \n\t" //load register 28 with data pointed at by X           //(r26,r27)  
    "increment register pair afterwards  
    "out __SP_L__, r28 \n\t" // restore stack pointer  
    "ld r29, x+ \n\t"  
    "out __SP_H__, r29 \n\t" // write contents of r29 to __SP_H__  
    "pop r31 \n\t" // pop registers  
    "pop r30 \n\t"  
    "....  
    "pop r1 \n\t"  
    "pop r0 \n\t" // restore status register before register r0  
    "out __SREG__, r0 \n\t"  
    "pop r0 \n\t"  
);
```



# Multitasking Operating Systems

# Overview

- Other services to be added to scheduler: timing, input/output (I/O) related features, semaphores
  - Example: Generating a time delay to lock out a task from running for some time, or waiting for an input
  - Consequence: Scheduler gets more complex
- Operating Systems (OS)
  - Subsumes scheduler and dispatcher
  - Provides abstraction of hardware
    - If OS runs on different hardware, application programmers do not need to know details of specific hardware
  - Goal: Reduction of ad hoc programming that must be done in every application
    - No need to reinventing the wheel in each application

# Preemptive Prioritized Multitasking OS

- Tasks are assigned priorities (high, medium, low, ...), and register with OS what events they are interested in
- Task is ready to run when one of the events it has registered for takes place
- Upon an event corresponding task is marked as ready to run
- If there is a task  $t_{hp}$  that is ready to run with higher priority than currently executing task  $t_r$ , then task  $t_r$  is preempted and control is transferred to  $t_{hp}$
- If running task runs out of things to do it yields voluntarily

# Preemptive Prioritized Multitasking OS

- Events are generated by ISRs
  - E.g. by I/O drivers supplied by OS
- ISRs can be part of OS or user supplied
- Events are stored up in FIFO queues, so if an event is not handled immediately it will not get lost
- Preemptive, prioritized systems impose a larger execution time and memory overhead than cooperative systems
  - Saving and restoring the context consumes quite an amount of memory and processor time



# Summary

- Superloops
  - Perfectly fine for multitasking applications of moderate complexity
  - Not an OS, no built in services
- Round robin scheduler
  - Powerful mechanism for structuring a program, provides services such as timers, input waits etc., may result in long delays for particular events
- Event Driven Loop
  - Order of tasks is determined dynamically
  - Difficult to analyze due to asynchronous events
- Cooperative multitasking OS
  - Simple, reliable, and provide a lot of useful functionality
  - OK for single programmer, but demands strict programming discipline
- Preemptive prioritized OS
  - Generally considered best for large programs
  - Less demanding of discipline from programmers than cooperative systems
  - Causes overhead



# Concurrent Task Model

# Overview

- In a multitasking OS several copies of a task may run simultaneously or at different times
- OS manages tasks
- Usage scenarios for simultaneously running tasks
  - Multiple rates: multimedia, automotive
  - Asynchronous input: user interfaces, communication systems
- Concurrent systems consist of a set of tasks

# Concurrent Task Model

- Difficult to write concurrent system using sequential program model (user has to program interleaving actions)
- **Concurrent Task Model**  
Describes functionality of system in terms of concurrently executing subtasks disregarding interleaving aspects
- Concurrent task model easier
  - Separate sequential programs (e.g. C function for each task)
  - Tasks communicate with each other (e.g. via shared variables)
- Systems which are inherently multitasking are easier to describe with concurrent task model

# Concurrent Task Model


```

ConcurrentTaskExample() {
  x = ReadX();
  y = ReadY();
  Call concurrently:
  PrintHelloWorld(x) and
  PrintHowAreYou(y);
}
PrintHelloWorld(x) {
  while(1) {
    print "Hello world";
    delay(x);
  }
}
PrintHowAreYou(y) {
  while(1) {
    print "How are you?";
    delay(y);
  }
}

```

Simple concurrent task example

- Simple example:
  - Read two integral numbers  $X$  and  $Y$
  - Display “Hello world.” every  $X$  seconds
  - Display “How are you?” every  $Y$  seconds
- More effort would be required with sequential program or state machine model
- Even better
  - OS takes care of timing!



Subroutine execution over time

```

Enter X: 1
Enter Y: 2
Hello world.    (Time = 1 s)
Hello world.    (Time = 2 s)
How are you?    (Time = 2 s)
Hello world.    (Time = 3 s)
How are you?    (Time = 4 s)
Hello world.    (Time = 4 s)
...

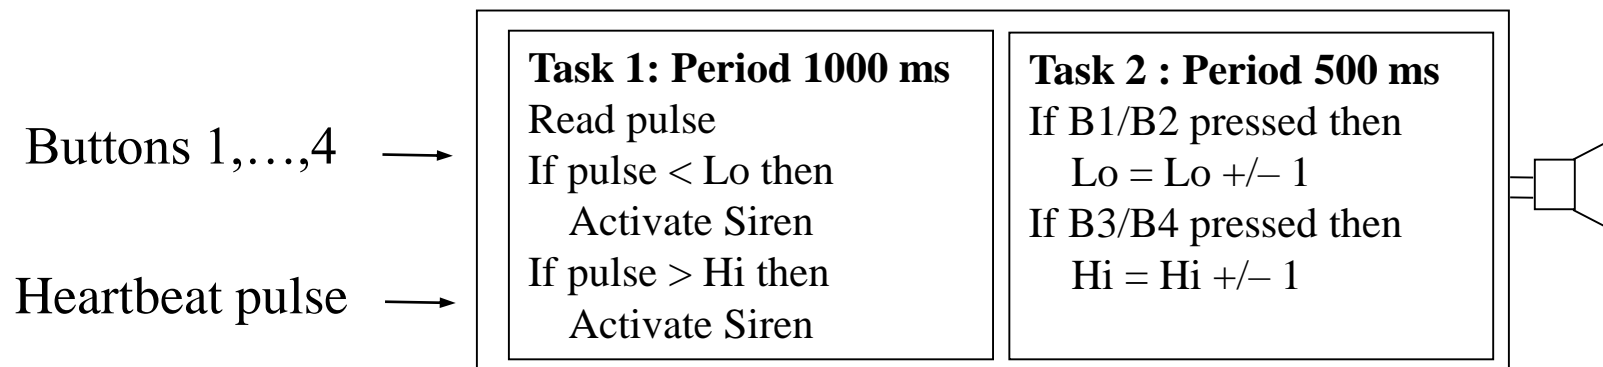
```

Sample input and output

- How to implement *Call concurrently*?

# Example I: Heartbeat Monitoring System

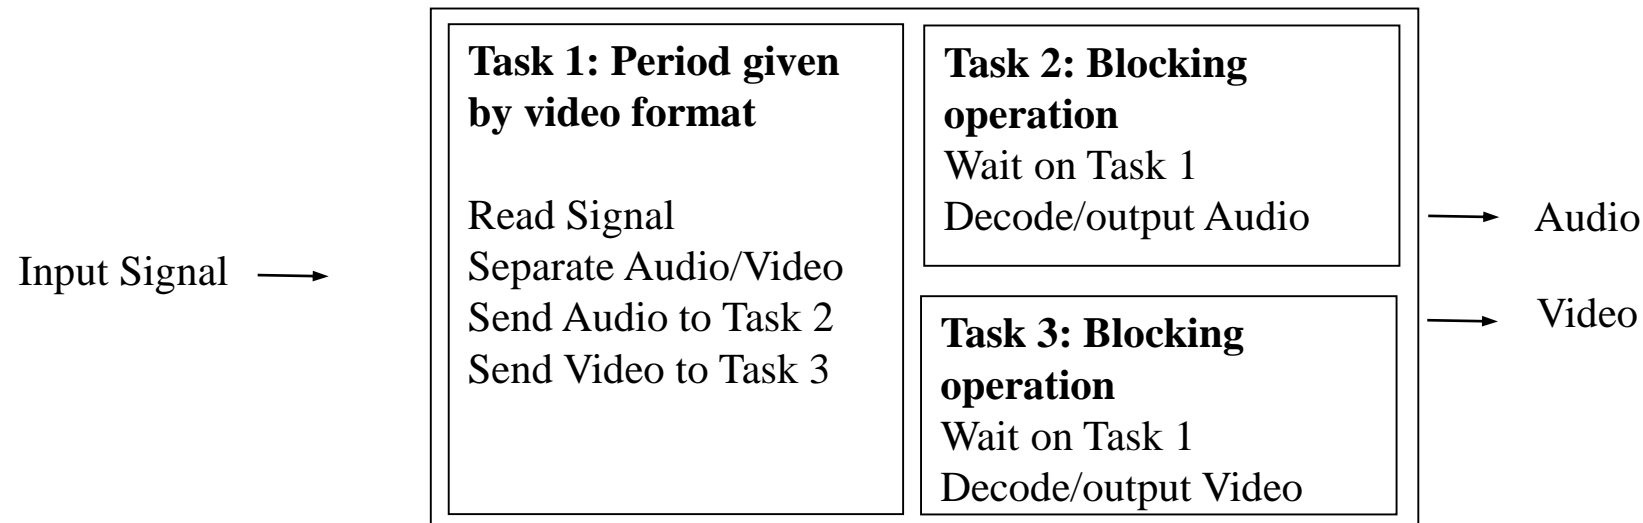
- Task 1:
  - Once per second sample input pulse and compute heartbeat of patient, signal if value is outside thresholds  $L_o$  and  $H_i$
- Task 2:
  - Twice a second check buttons: if one of four buttons is pressed, increment/decrement  $L_o$  resp.  $H_i$
- Tasks are independent, but access common data ( $L_o$ ,  $H_i$ )
- OS schedules tasks according to periods, user not responsible for timing and interleaving the code





# Example II: Set-top Box

- Task 1:
  - Receive signal from antenna, decompose into compressed video and audio
- Tasks 2 & 3:
  - Decode compressed video resp. audio
- Tasks are independent, but access common data
- OS schedules tasks (blocking and unblocking)



# Basic Operations on Tasks

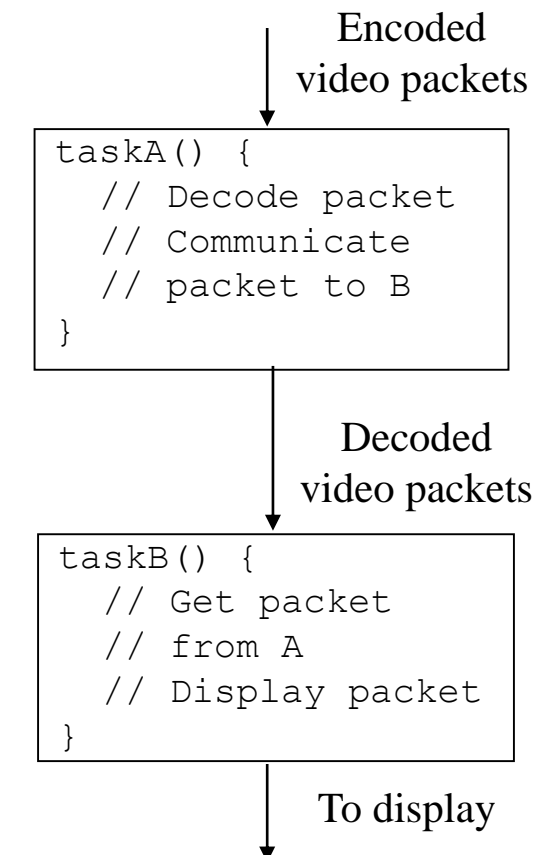

- Create and terminate
  - Create is like a procedure call but caller doesn't wait (block)
    - Created task can itself create new tasks
    - In HelloWorld/HowAreYou example, two tasks are created
  - Terminate kills a task, destroying all data
- Suspend and resume
  - Suspend puts task on hold, saving state for later execution
  - Resume starts the task again where it left off
- Join
  - Task suspends until a particular child task finishes execution
- Wait (block)
  - Task waits for condition to be fulfilled by other task (avoiding busy waiting)



# Communication Among Tasks

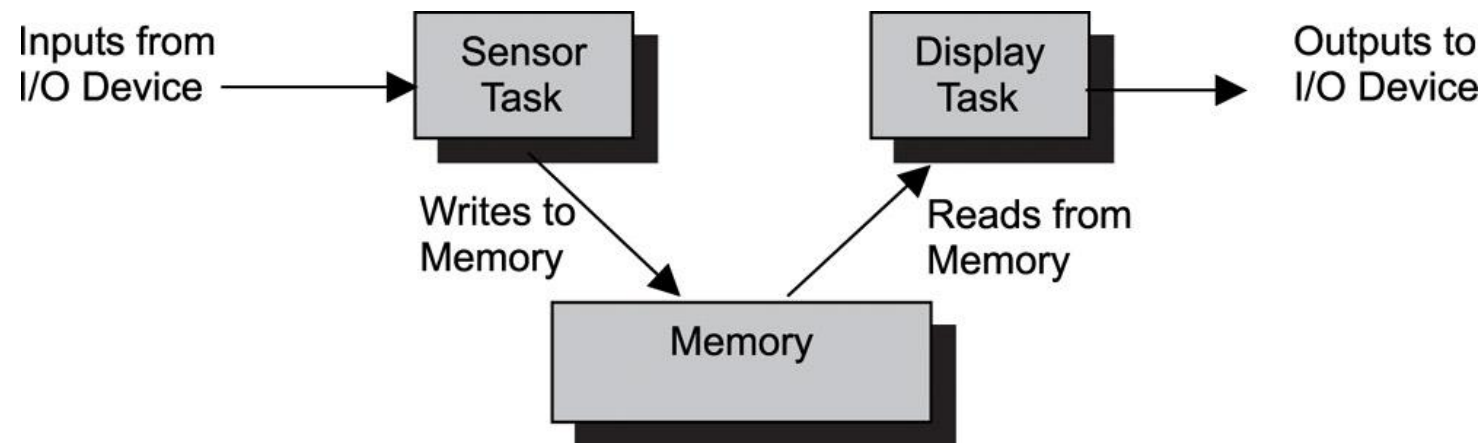
# 4. Communication Among Tasks

- Tasks need to communicate data and signals to solve their computation problem
  - Tasks that don't communicate are just independent programs solving separate problems
- Basic pattern: producer/consumer
  - Task A produces data and B consumes it
  - Example: A decodes video packets, B displays decoded packets on a screen
- How do we achieve this communication?
  - Two basic methods
    - Shared memory
    - Message passing
  - What to do if tasks run on variable speeds?
    - Buffers



# Shared Memory

- Tasks read/write shared variables
  - No time overhead, easy to implement
  - Example with two tasks:
    - Sensor task periodically receives data from a sensor and writes data to shared memory
    - Display task periodically reads from shared memory and sends data to display
  - Access by multiple tasks **must** be synchronized to maintain the integrity of shared data





# Shared Memory - Example

## ■ Producer/consumer with preemptive scheduler

- Shared variables: *buffer[N]*, *count*
  - *count* = # of valid data items in *buffer*
- *taskA* produces data items & stores them in *buffer*
  - If *buffer* is full, then *taskA* must wait
- *taskB* consumes data items from *buffer*
  - If *buffer* is empty, then *taskB* must wait
- Race condition:  
Error when both tasks try to update *count* concurrently (lines 10 and 19) and the following execution sequence occurs: (let *count* = 3)
  - A loads *count* (3) from memory into register R1 (R1 = 3)
  - A increments R1 (R1 = 4)
  - B loads *count* (3) from memory into register R2 (R2 = 3)
  - B decrements R2 (R2 = 2)
  - A stores R1 back to *count* in memory (*count* = 4)
  - B stores R2 back to *count* in memory (*count* = 2)*count* now has incorrect value of 2
- No busy waiting in lines 7, 16

```
01: data_type buffer[N]; // shared
02: int count = 0; // shared

03: void taskA() {
04:     int i = 0; // next slot to store
05:     while(1) {
06:         produce(&data);
07:         while (count == N)
08:             wait_task;
09:         buffer[i] = data;
10:         i = (i + 1) % N;
11:         count = count + 1;
12:     }

13: void taskB() {
14:     int i = 0; // next slot to read
15:     while(1) {
16:         while (count == 0)
17:             wait_task;
18:         data = buffer[i];
19:         i = (i + 1) % N;
20:         count = count - 1;
21:         consume(&data);
22:     }

23: void main() {
24:     create_task(taskA);
25:     create_task(taskB);
26: }
```



# Message Passing

- Message passing
  - Data explicitly sent from one task to another
    - Sending task performs special operation *send* to *send data*
    - Receiving task performs special operation *receive* to receive the data
    - Both operations must explicitly specify which task it is sending to or receiving from
    - Receive is blocking, send is not blocking
    - Blocked task waits for another task to send data
  - Safer model, but less flexible

```
taskA() {  
    while(1) {  
        produce(&data)  
        send(B, &data);  
        /* do something */  
        receive(B, &data);  
        consume(&data);  
    }  
}
```


```
taskB() {  
    while(1) {  
        receive(A, &data);  
        transform(&data)  
        send(A, &data);  
        /* do something */  
    }  
}
```

# Back to Shared Memory: Mutual Exclusion

- Certain sections of code should not be performed concurrently
  - **Critical section:**  
Section of code where simultaneous updates, by multiple tasks to a shared memory location, can occur
- When a task enters the critical section, all other tasks must be locked out until it leaves the critical section
- Solution with Mutex (requires hardware support)
  - A shared object used for locking and unlocking segment of shared data
  - Disallows read/write access to memory it guards
  - Multiple tasks can perform lock operation simultaneously, but only **one** task will acquire lock
  - All other tasks trying to obtain lock are put in blocked state until unlock operation performed by acquiring task when it exits critical section
  - These tasks will then compete for lock again



# Mutex



# Implementing a Lock

- Shared Boolean variable: `locked`
- Value indicates whether access to critical section is currently allowed

- Access to variable `locked` using two functions

```
void lock() { //blocking function
    .....
}
void unlock() {
    .....
}
```

- Usage: Placement of code of critical section

```
lock();
    // code of critical section
unlock();
```

# Implementing a Lock

- locked: pointer to lock variable

```
void lock() {  
    while (*locked);  
    *locked = true;  
}
```



- Atomic operations on some processors
  - test-and-set, fetch-and-add, compare-and-swap
- **Logical** implementation of test-and-set (in reality one machine instruction)

```
boolean testAndSet(boolean * target) {  
    boolean rv = * target;  
    * target = true;  
    return rv;  
}
```

- Usage

```
void lock() {  
    while (testAndSet(locked)); //busy waiting  
}  
void unlock() {  
    locked = false;  
}
```

# Consumer-Producer Problem with Mutex

- *mutex* is used to ensure that critical sections are executed in mutual exclusion of each other
- Same execution sequence as before:
  - *A/B* execute *lock* operation on *count\_mutex*
  - Either *A* or *B* will acquire *lock*
    - Say *B* acquires it
    - *A* will be put in blocked state
  - *B* loads *count* (3) from memory into register R2 ( $R2 = 3$ )
  - *B* decrements R2 ( $R2 = 2$ )
  - *B* stores R2 back to *count* in memory (2)
  - *B* executes *unlock* operation
    - *A* is placed in runnable state again and acquires lock
  - *A* loads *count* (2) from memory into register R1 ( $R1 = 2$ )
  - .....
- *Count* now has correct value of 3

```
01: data_type buffer[N];
02: int count = 0;
03: mutex count_mutex;

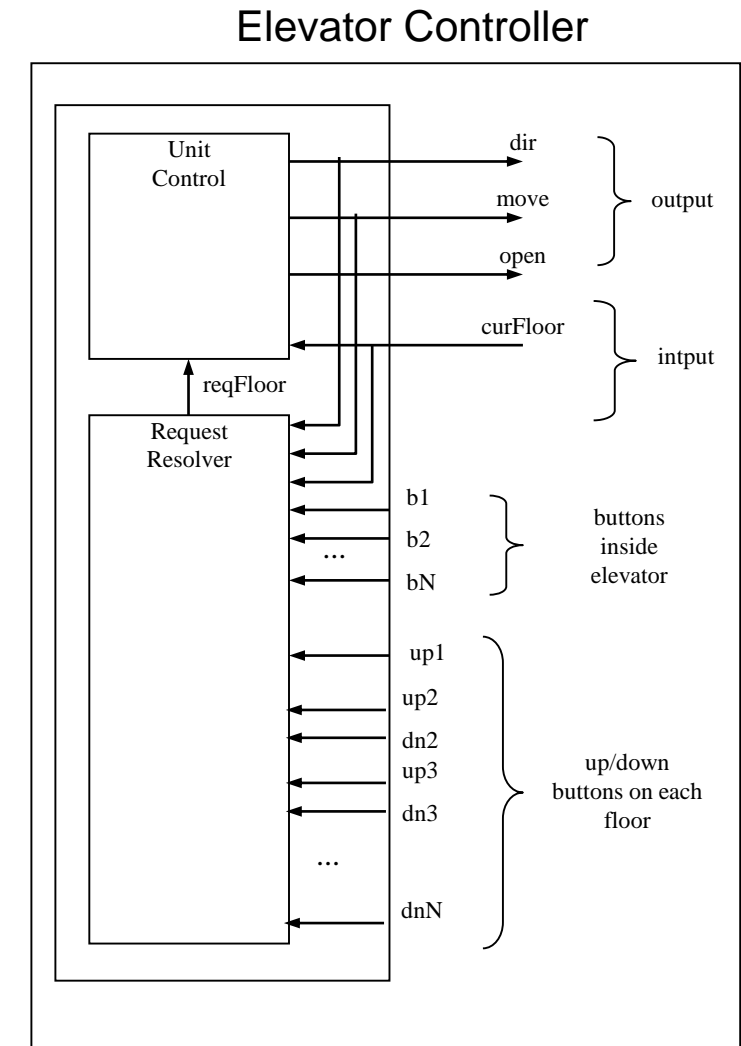
04: void taskA() {
05:     int i = 0;
06:     while(1) {
07:         produce(&data);
08:         while (count == N)
09:             wait_task;
10:         buffer[i] = data;
11:         i = (i + 1) % N;
12:         count_mutex.lock();
13:         count = count + 1;
14:         count_mutex.unlock();
15:     }

16: void taskB() {
17:     int i = 0;
18:     while(1) {
19:         while (count == 0)
20:             wait_task;
21:         data = buffer[i];
22:         i = (i + 1) % N;
23:         count_mutex.lock();
24:         count = count - 1;
25:         count_mutex.unlock();
26:         consume(&data);
27:     }
```



# Example: Task Communication

- Elevator controller
- Unit control task:
  - reads next requested floor from queue and moves to requested floor
- Request resolver task:
  - appends new requests to queue
  - Consumer-Producer problem
- Realization:
  - Shared variable for queue
  - Mutex for controlling access to queue



# Mutual Exclusion with AVR

- `#define ATOMIC_BLOCK (type)`
  - Creates block of code guaranteed to be executed atomically
  - Upon entering block Global Interrupt Status flag in SREG is disabled
  - State of SREG upon exit is controlled by parameter `type`
    - `#define ATOMIC_FORCEON`
      - causes ATOMIC\_BLOCK to force state of SREG on exit, enabling the Global Interrupt Status flag bit (i.e. previous value of SREG register is not need saved at start of block)
    - `#define ATOMIC_RESTORESTATE`
      - causes ATOMIC\_BLOCK to restore previous state of SREG register, saved before Global Interrupt Status flag bit was disabled

# Deadlocks in Concurrent Programming

- Deadlock: Condition where two tasks are blocked waiting for other to unlock critical sections
  - Both tasks are then in blocked state and
  - cannot execute unlock operation and wait forever
- Example: Two different critical sections that can be accessed simultaneously
  - Two locks needed (*mutex1*, *mutex2*)
  - Following execution sequence produces deadlock
    - *A* executes lock operation on *mutex1* (and acquires it)
    - *B* executes lock operation on *mutex2* (and acquires it)
    - *A/B* both execute in critical sections 1 and 2, resp.
    - *A* executes lock operation on *mutex2*
      - *A* blocked until *B* unlocks *mutex2*
    - *B* executes lock operation on *mutex1*
      - *B* blocked until *A* unlocks *mutex1*
    - DEADLOCK!
- Possible deadlock elimination protocols
  - Acquire all needed locks in one step
  - Number locks and acquire them in increasing order

```
01: mutex mutex1, mutex2;

02: void taskA() {
03:     while(1) {
04:         ...
05:         mutex1.lock();
06:         /* critical section 1 */
07:         mutex2.lock();
08:         /* critical section 2 */
09:         mutex2.unlock();
10:         /* critical section 1 */
11:         mutex1.unlock();
12:     }
13: }

14: void taskB() {
15:     while(1) {
16:         ...
17:         mutex2.lock();
18:         /* critical section 2 */
19:         mutex1.lock();
20:         /* critical section 1 */
21:         mutex1.unlock();
22:         /* critical section 2 */
23:         mutex2.unlock();
24:     }
25: }
```

# Synchronization Among Tasks

- Concurrently running tasks must synchronize their execution when a task must wait for
  - another task to compute some value
  - reach a known point in their execution
  - signal some condition
- Recall producer-consumer problem
  - *taskA* must wait if *buffer* is full
  - *taskB* must wait if *buffer* is empty
- Busy-waiting
  - Tasks execute loops instead of being blocked
  - CPU time is wasted!
- More efficient methods avoid busy waiting
  - Join operation, and blocking send and receive discussed earlier
    - Both block task, no wastage of CPU time
  - Condition variables and monitors avoid busy waiting



# SES

## Chapter 6: Operating Systems for Embedded Systems

Prof. Dr.-Ing. Bernd-Christian Renner

