

SES

Chapter 10: Real-Time Embedded Systems

Prof. Dr.-Ing. Bernd-Christian Renner



Contents

1. Definitions and Examples
2. Soft and Hard Real Time
3. Basic Notions
4. Scheduling
5. Specific Scheduling Algorithms
 - Aperiodic tasks
 - Periodic tasks
6. Real-time Operating Systems



Definitions and Examples


What is a Real-Time System?

- Real-time system
 - Information processing activity which has to respond to externally generated input stimuli within a specified period otherwise risks severe consequences, including failure
- Logical correctness of real-time system is based on
 - correctness of outputs and
 - their timelines

Related Notions

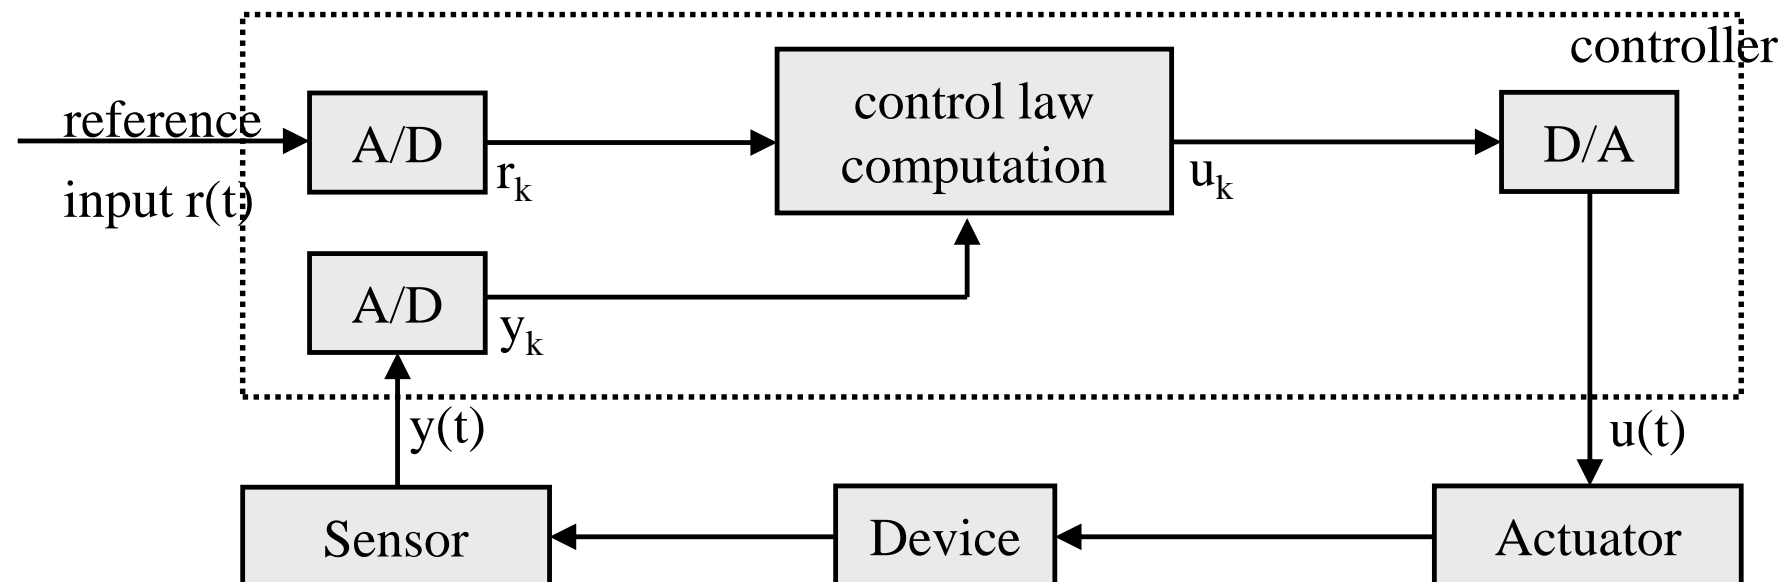
- Reactive system
 - Continuous interaction with the environment (as opposed to information processing)
- Embedded system
 - Computer system encapsulated in its environment (device it controls), combination of computer hardware and software, dedicated to specific purpose
- Safety-critical system
 - A failure may cause injury, loss of lives, significant financial loss

Block Diagram of generic Real-Time System



Example 1: Digital Process Control

- Controlling device with actuator, based on sampled sensor data
 - $y(t)$: measured state of device
 - $r(t)$: desired state of the device
 - Calculate control output $u(t)$ as a function of $y(t)$, $r(t)$



Example 1: Digital Process Control

- Pseudo-code for controller:

```
set timer to interrupt periodically with period T;  
at timer interrupt, do  
    analogue-to-digital conversion of  $y(t)$  to get  $y_k$ ;  
    analogue-to-digital conversion of  $r(t)$  to get  $r_k$ ;  
    compute control output  $u_k$  based on reference  $r_k$  and  $y_k$ ;  
    digital-to-analogue conversion of  $u_k$  to get  $u(t)$ ;  
end do;
```

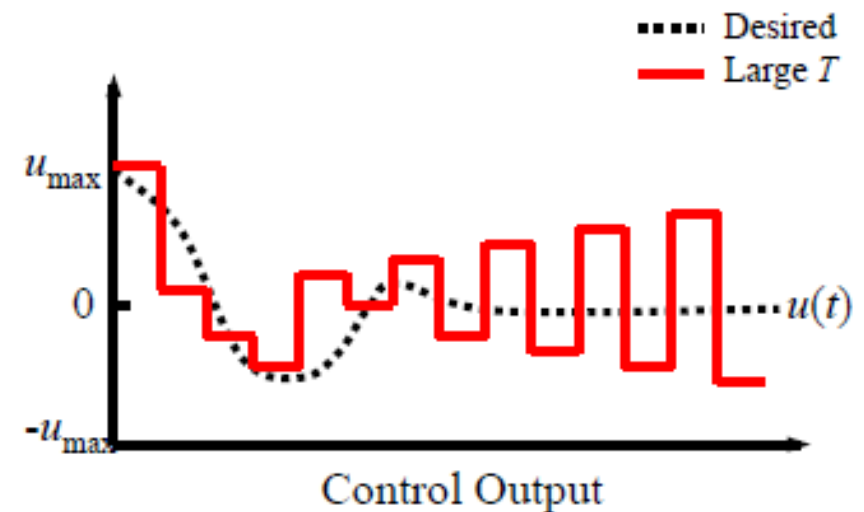
- Sampling period T

- time between consecutive measurements of $y(t)$, $r(t)$

- Hardware must **guarantee** that all conversions and processing is possible in time T
- How can this be guaranteed?


Example 1: Digital Process Control

- Effective control of device depends on
 - Correct control law computation and reference input
 - Accuracy of the sensor measurements:
 - Resolution of sampled data (i.e. bits per sample)
 - Length of T (i.e. samples per second, $1/T$)
- Small T better approximates analogue behavior
 - Small T requires
 - more processor-time
 - better ADC hardware
 - Downside:
 - possibly lower resolution
- Large T results in oscillation



Multi-rate systems

- Tasks may be synchronous or asynchronous
- Synchronous tasks may recur at different rates
- Processes run at different rates based on computational needs of the tasks
- Example 2: Engine control
 - spark control
 - crankshaft & oxygen sensing
 - fuel/air mixture
 - multi mode operation:
warm-up, cruise, climbing steep hills



Typical Rates in Engine Controllers

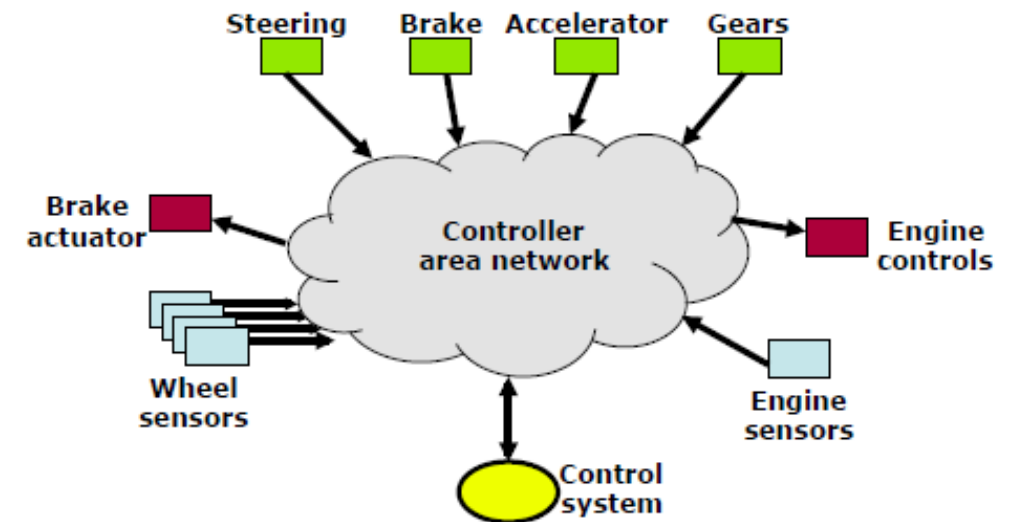
Variable	Update period (ms)
Engine spark timing	2
Throttle (Gaspedal)	2
Air flow	4
Battery voltage	4
Fuel flow	10
Recycled exhaust gas	25
Status switches	20
Air temperature	400
Barometric pressure	1000
Spark (dwell)	1
Fuel adjustment	8
Carburetor (Vergaser)	25

Real-Time Communications

- Real-time systems are increasingly distributed, including communication networks
 - Control loop may include a communication step
 - System may depend on network stimuli
- Not only does a system need to run a control law with time constraints, it must also
 - schedule communications
 - send and receive messages according to deadlines

Example 3: Drive by Wire

- All data must be delivered reliably
 - Bad if you turn steering wheel, and nothing happens
- Systems are prioritized
 - Anti-lock brakes have a faster response time than driver, so prioritize to ensure car doesn't skid
 - Commands from control system have highest priority, then sensors and actuators
- Network must schedule and prioritize communications





Soft and Hard Real Time

Soft and Hard Real Time

■ Hard RT System

- missing a deadline may cause failure of system, e.g., aircraft control, nuclear plant control
 - Side airbag in car, reaction in <10 ms
 - Wing vibration of airplane, sensing every 5 ms

■ Soft RT System

- meeting a deadline is highly desirable for performance reasons, e.g., multimedia application, booking system, displaying status information

Real time vs best effort:

- best effort = low average time
- real time = predictability, bounded worst case time

Soft and Hard Real Time

- Most systems contain both hard and soft deadlines
- Essential for hard real time: Upper bounds of execution times of all tasks must be known at compile time
 - Commonly called **Worst-Case Execution Time** (WCET)
- Further notion: Firm deadline
 - Missing a deadline makes task useless (similar to hard deadline), however deadline may be missed occasionally (similar to soft deadline)
- Tasks may have cost functions associated with them for missing their deadline
 - System tries to minimize costs

Predictability

- Predictability: Correctness of prediction of system's state
- Predictability
 - is one of the most important requirements
 - but also one of the most difficult requirements to achieve, in particular in modern processors:
 - cache, pipelines, branch prediction, interrupt handling
 - memory management
 - priority inversion
 - difficult to calculate WCET
 - ...

Types of Real-Time Applications

- Purely periodic
 - Every task executes periodically
 - Demands in (computing, communication, and storage) resources do not vary significantly from period to period
 - Example
 - Most digital controllers and real-time monitors
- Mostly periodic
 - Most tasks execute periodically
 - System must also respond to some external events (fault recovery and external commands) asynchronously
 - Example
 - Modern avionics and process control systems

Types of Real-Time Applications

- Asynchronous: Mostly predictable
 - Most tasks are not periodic
 - Time between consecutive executions of a task may vary, but variance is bounded
- Asynchronous: Unpredictable
 - Applications that react to asynchronous events and have tasks with unpredictable high variance in run-time
 - Variations have neither bounded ranges nor known statistics



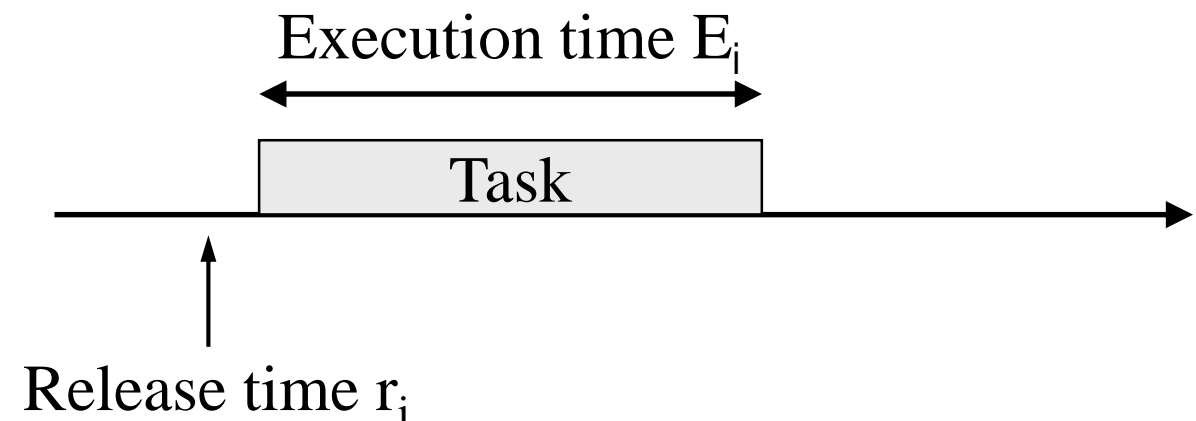
Basic Notions

Basic Notions

- **Task** (a.k.a. process, job)
 - Something that needs to be done (sequentially)
- Given a set of tasks $J = \{J_1, J_2, \dots, J_n\}$
- **Schedule**
 - Assignment of execution times for tasks such that all tasks are executed until completion
- Formal definition of schedule σ :
 - $\sigma: T \rightarrow J$ where $\sigma(t)$ denotes task which is executed at time t
 - If $\sigma(t) = 0$ then processor is called **idle**
- **Context switch**
 - Times when σ changes its value
- **Preemptive schedule**
 - Schedule in which running task can be suspended at any time

Basic Notions

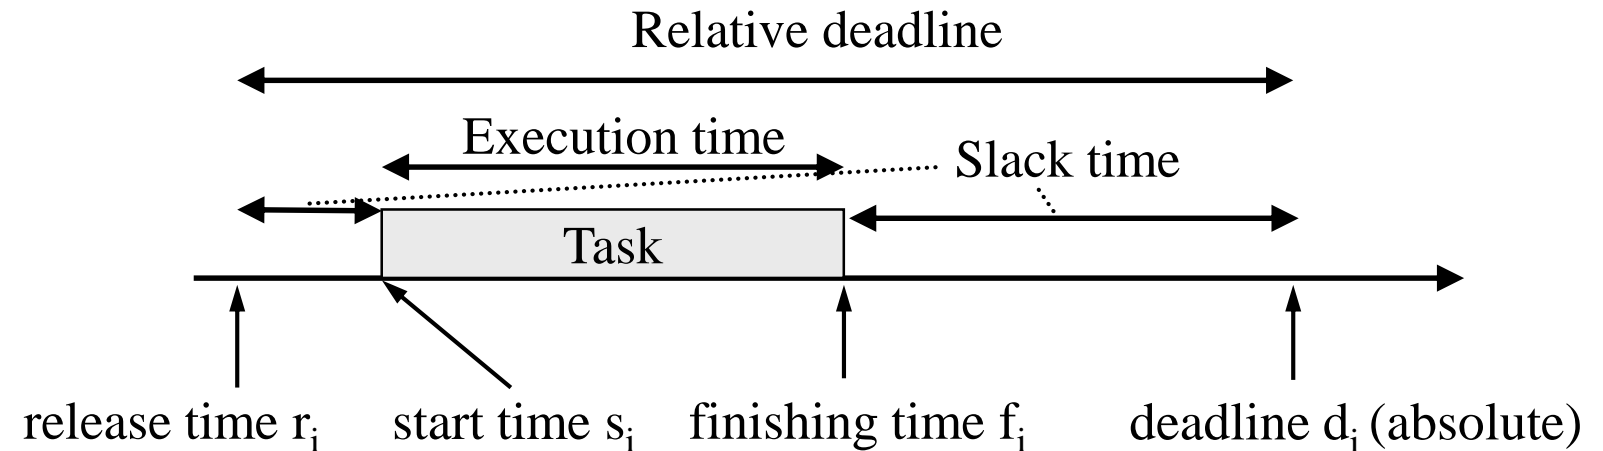
- A schedule is said to be **feasible**, if all tasks can be completed according to a set of specified constraints, e.g., deadlines
- A set of tasks is said to be **schedulable**, if there exists at least one feasible schedule
- **Release time** (a.k.a. arrival time) r_i
 - time at which a task becomes ready for execution
- **Execution time** E_i
 - time necessary for executing task without interruption



Task Execution Characteristics

- Execution time E_i
 - Execution time in absence of preemption (seconds or clock cycles)
 - Problem: Execution time varies (caching, pipelining,..)
 - Real-time systems: Only WCET matters!
- General schedulability analysis is NP-hard problem
 - For aperiodic tasks EDF (explained later) runs in polynomial time


Basic Notions




- **Deadline d_i**
 - time at which a task must be completed
- **Start time s_i**
 - time at which a task starts its execution
- **Finishing time f_i**
 - time at which a task finishes its execution
- **Slack time $X_i = d_i - r_i - E_i$**
 - Maximum time a task can be delayed on its activation to complete within its deadline

Basic Notions

- $d_i \geq r_i + E_i$
- **Lateness:** $L_i = f_i - d_i$
 - represents delay of a task completion with respect to its deadline
 - if task completes before deadline, lateness is negative




Example



- Execution times: $E_1 = 9$, $E_2 = 12$
- Release times: $r_1 = 0$, $r_2 = 4$
- Deadlines: $d_1 = 22$, $d_2 = 27$ ► Slack time: $X_1 = 13$, $X_2 = 11$
- Start times: $s_1 = 0$, $s_2 = 6$
- Finishing times: $f_1 = 18$, $f_2 = 28$ ► Lateness: $L_1 = -4$, $L_2 = 1$


Periodic tasks

- Periodic task:
 - Sequence of identical activities (same execution time) that are regularly activated at a constant rate with period P
 - Release time of first instance is called **phase** Φ
 - Release time of k^{th} activation is $\Phi + (k - 1)P$
 - Deadline of k^{th} activation is $d + (k - 1)P$ (often $d = P$)



Example: Heating Furnace

- System takes 20 ms to initialize after turned on
- After initialization, every 100 ms, the system:
 - samples and reads temperature sensor
 - computes control-law for furnace to process temperature readings
 - determines correct flow rates of fuel, air and coolant
 - adjusts flow rates to match computed values
- Release time of k^{th} activation is $20 + (k - 1) \times 100$ ms






Scheduling

Scheduling Goals

- Produce feasible schedules
- Minimize given cost function
- Metrics
 - Average response time
 - Total completion time
 - For soft RT systems
 - Minimize maximum lateness & number of late tasks
- Scheduling process consumes CPU time
 - Thus, not all CPU time is available for tasks
 - Scheduling overhead must be taken into account for exact schedule
 - Assumption in the following:
 - All overheads caused by scheduler are assumed to be zero

Scheduling and Context Switch Overhead

Task	Execution time E_i	Period = Deadline d_i
J_1	3	10
J_2	3	5



- No feasible schedule exists if context switch overhead is 1
 - $E_1 + 2E_2 + 2 = 11$, hence earliest finishing time is 12
 - Second invocation of J_2 must be finished before 11

Scheduling Algorithms

- Static algorithms
 - Scheduling decisions are based on fixed parameters, assigned to tasks before their activation
- Dynamic algorithms
 - Scheduling decisions are based on dynamic parameters that may change during execution

Scheduling Based on Priorities

- **Principle of priority scheduling**
 - Upon each scheduling event (task finishes, task released) task with highest priority in task set is scheduled, if currently executing task has lower priority then it is interrupted
 - Ties are broken arbitrarily, e.g. by FIFO
- Task set = released but not completed tasks
- Task set is dynamically maintained, implemented as priority queue
- Priorities are determined
 - statically during creation or
 - dynamically during execution
- All scheduling algorithms are based on priorities



5

Specific Scheduling Algorithms

Specific Scheduling Algorithms

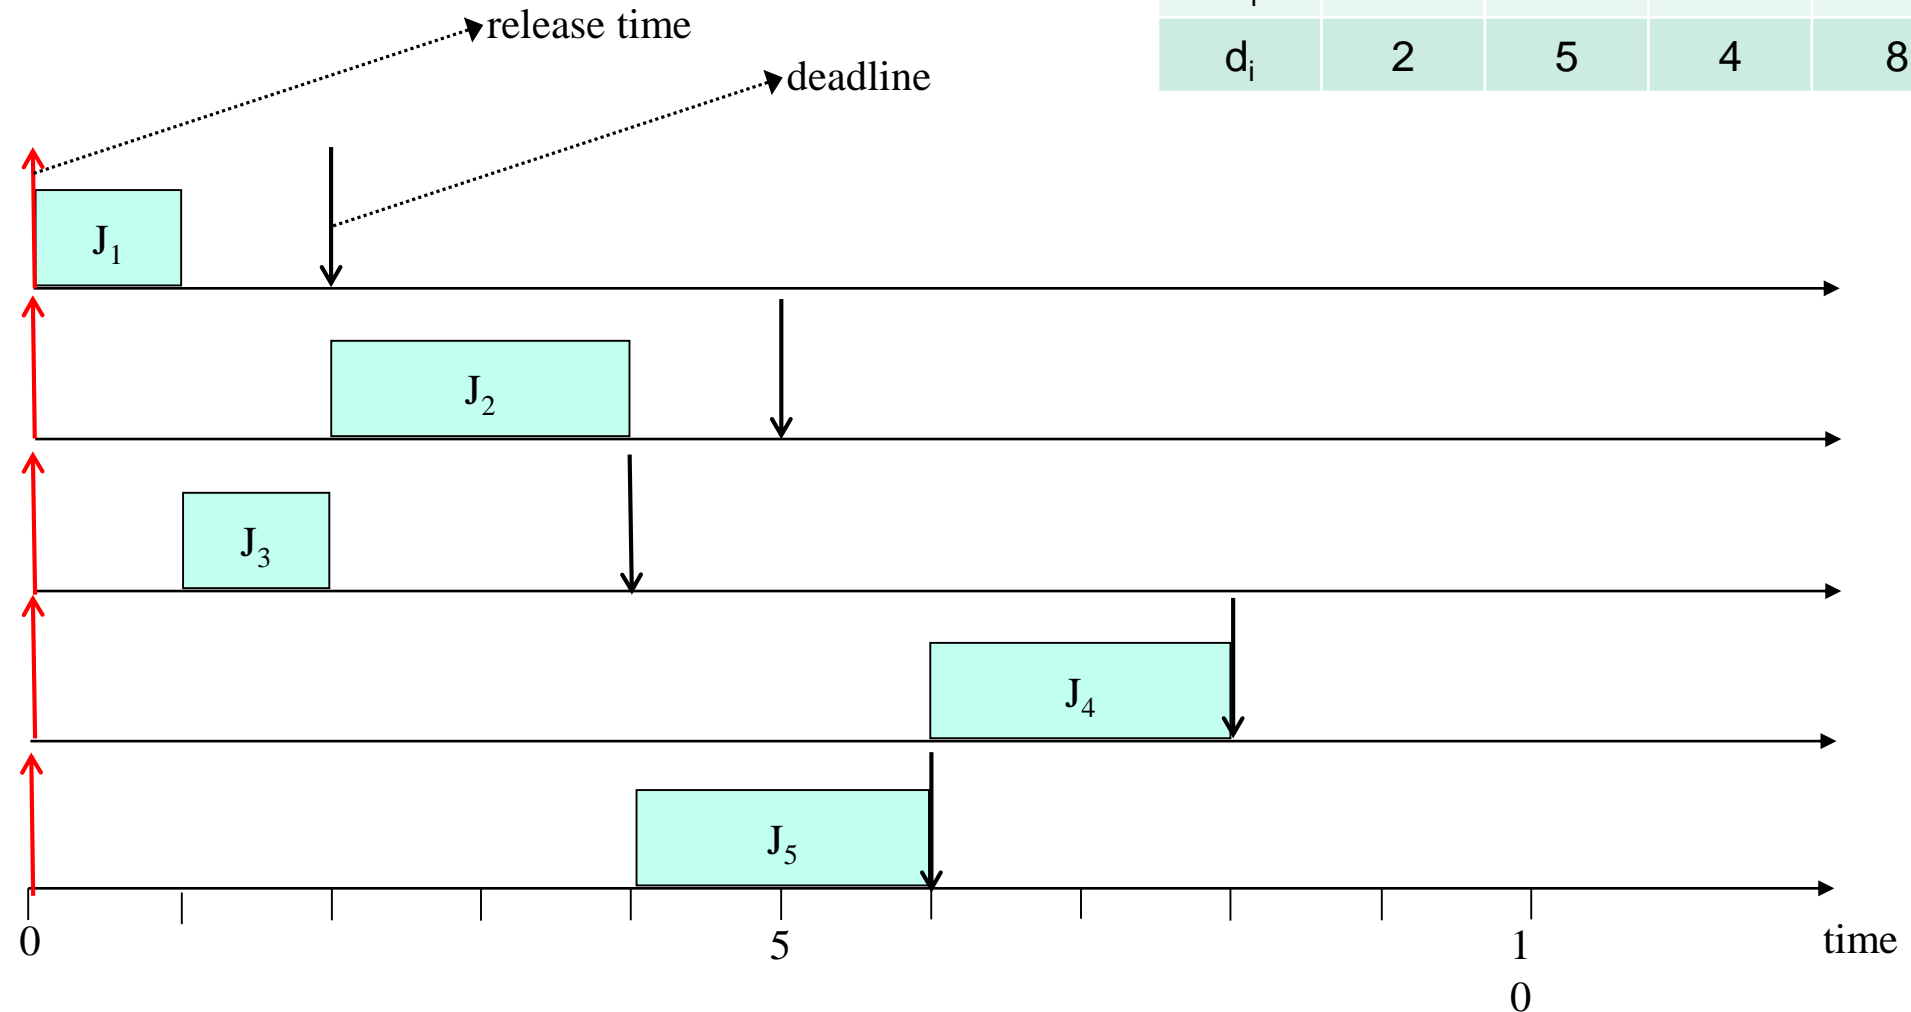
- Aperiodic tasks
 - Earliest Deadline First (EDF)
 - preemptive schedule
 - arbitrary arrival times
 - independent tasks (i.e. no precedence constraints)
 - Earliest Deadline First (EDF*)
 - With precedence constraints
- Periodic tasks
 - Rate Monotonic Scheduling (RM)
 - no precedence constraints
 - deadlines equal periods
 - Deadline Monotonic Scheduling (DM)
 - Same as RM but deadlines can be different from periods

Earliest Deadline First (EDF)

- Priorities are based on deadlines
 - Tasks with earlier deadlines have higher priorities
- Guarantees
 - If a set of tasks has a feasible schedule, then EDF will schedule these tasks, so they all complete by their deadline
 - Schedule is optimal with respect to minimizing maximum lateness
- EDF can also be used for periodic tasks


EDF: Example 1

	J_1	J_2	J_3	J_4	J_5
r_i	0	0	0	0	0
E_i	1	2	1	2	2
d_i	2	5	4	8	6



EDF: Example 2

	J ₁	J ₂	J ₃	J ₄	J ₅
r _i	0	0	2	3	6
E _i	1	2	2	2	2
d _i	2	5	4	10	9



Concept of proof for guarantee of EDF

- For each time interval $[t, t+1)$ it is verified, whether actual running task is the one with earliest absolute deadline
- If this is not the case, task with earliest absolute deadline is executed in this interval instead
- This operation cannot increase the maximum lateness

Schedulability test for EDF

- Assume there exist a feasible schedule for a set of n tasks
- Consider a new task arriving at time t (i.e. now $n+1$ tasks)
- Order all tasks by deadline $J = \{J_1, J_2, \dots, J_{n+1}\}$
- Denote by $c_i(t)$ the remaining worst-case execution time of task J_i at time t for this schedule
- - $f_0 := t$
 - for** $i := 1, \dots, n+1$
 - $f_i := f_{i-1} + c_i(t)$
 - if** ($f_i > d_i$)
 - return** *impossible to schedule*
 - return** *possible to schedule*

Pros and Cons of EDF

- Pros
 - Simple and works nicely in theory
 - Simple schedulability test
 - Optimal
 - Best CPU utilization
- Cons
 - Difficult to implement in practice. Not very often adopted due to dynamic priority-assignment (maintenance of ready queue is expensive)
 - Non stable: if any task instance fails to meet its deadline, the system is not predictable, any instance of any task may fail

EDF*: EDF With Precedence Constraints

- In some applications tasks have precedence constraints
 - A task must be completed before start of some other tasks
- Requirements
 - In a valid schedule a task starts execution
 - not earlier than its release time and
 - not earlier than the finishing times of its predecessors
 - All tasks finish their execution within their deadlines
- EDF* algorithm
 - Determines a feasible schedule for tasks with precedence constraints if there exists one

Principle of EDF*

1. Transform set J of dependent tasks into set J^* of independent tasks with new release times r^* and new deadlines d^*
 2. Apply EDF to set J^*
- Transformation is done by modification of
 - Release times
 - Task must not start execution earlier than minimum finishing time of its predecessors and its own release time
 - Deadlines
 - Task must finish execution time within its deadline
 - Task must not finish execution later than maximum start time of its successor

EDF*

Algorithm for modification of release times:


1. For any initial node of precedence graph set $r_i^* = r_i$
2. Select a task J_k such that its release time has not been modified, but release times of all immediate predecessors J_i have been modified
3. If no such task exists, halt
else
set $r_k^* = \max \{r_k, \max\{r_i^* + E_i : J_i \rightarrow J_k\}\}$
4. Return to second step

EDF*


Algorithm for modification of deadlines:

1. For all terminal nodes of precedence graph set $d_i^* = d_i$
2. Select a task J_k such that its deadline has not been modified but deadline of all immediate successors J_i have been modified
3. If no such task exists, halt
else
set $d_k^* = \min \{d_k, \min\{d_i^* - E_i : J_k \rightarrow J_i\}\}$
4. Return to second step

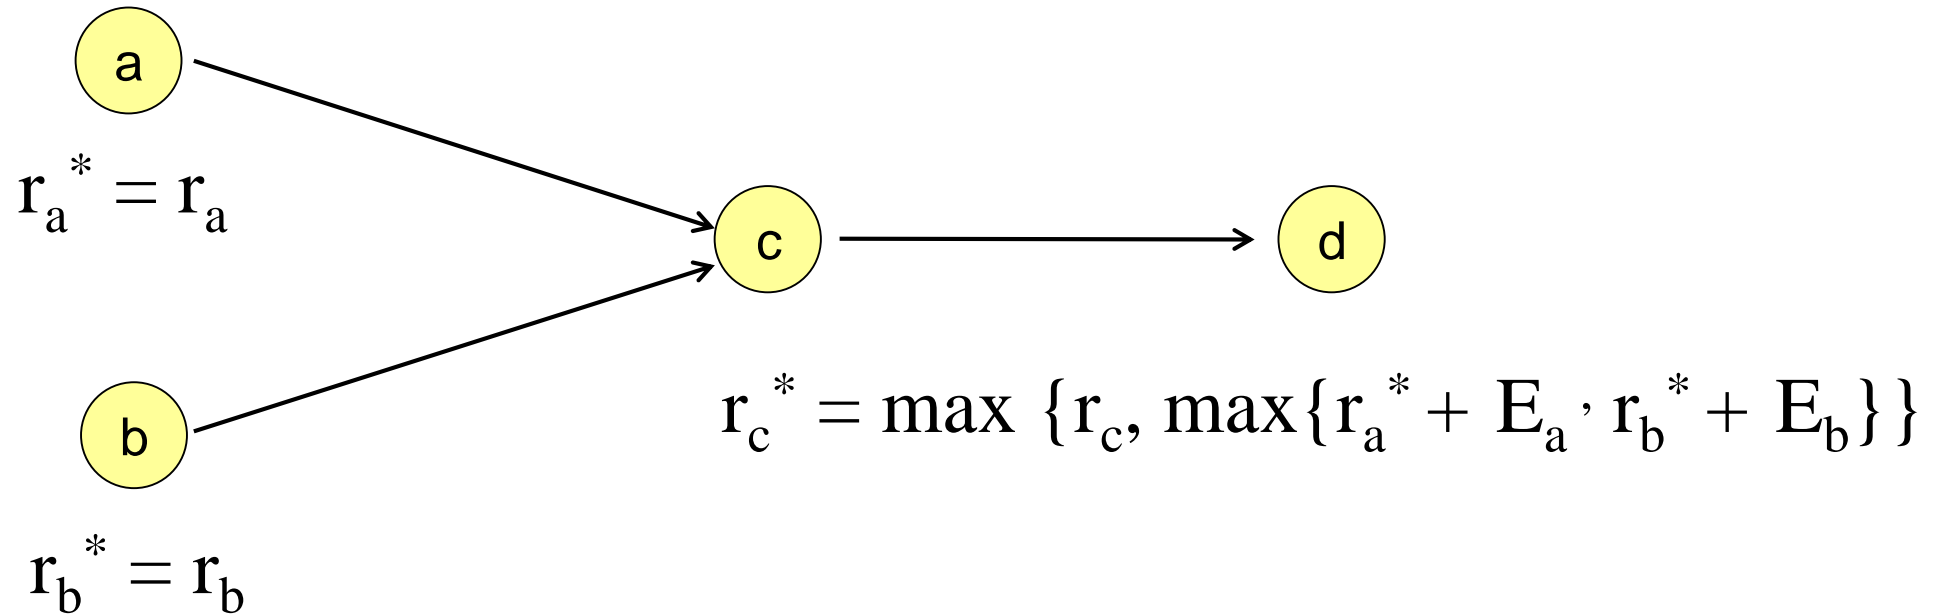
Example: EDF*



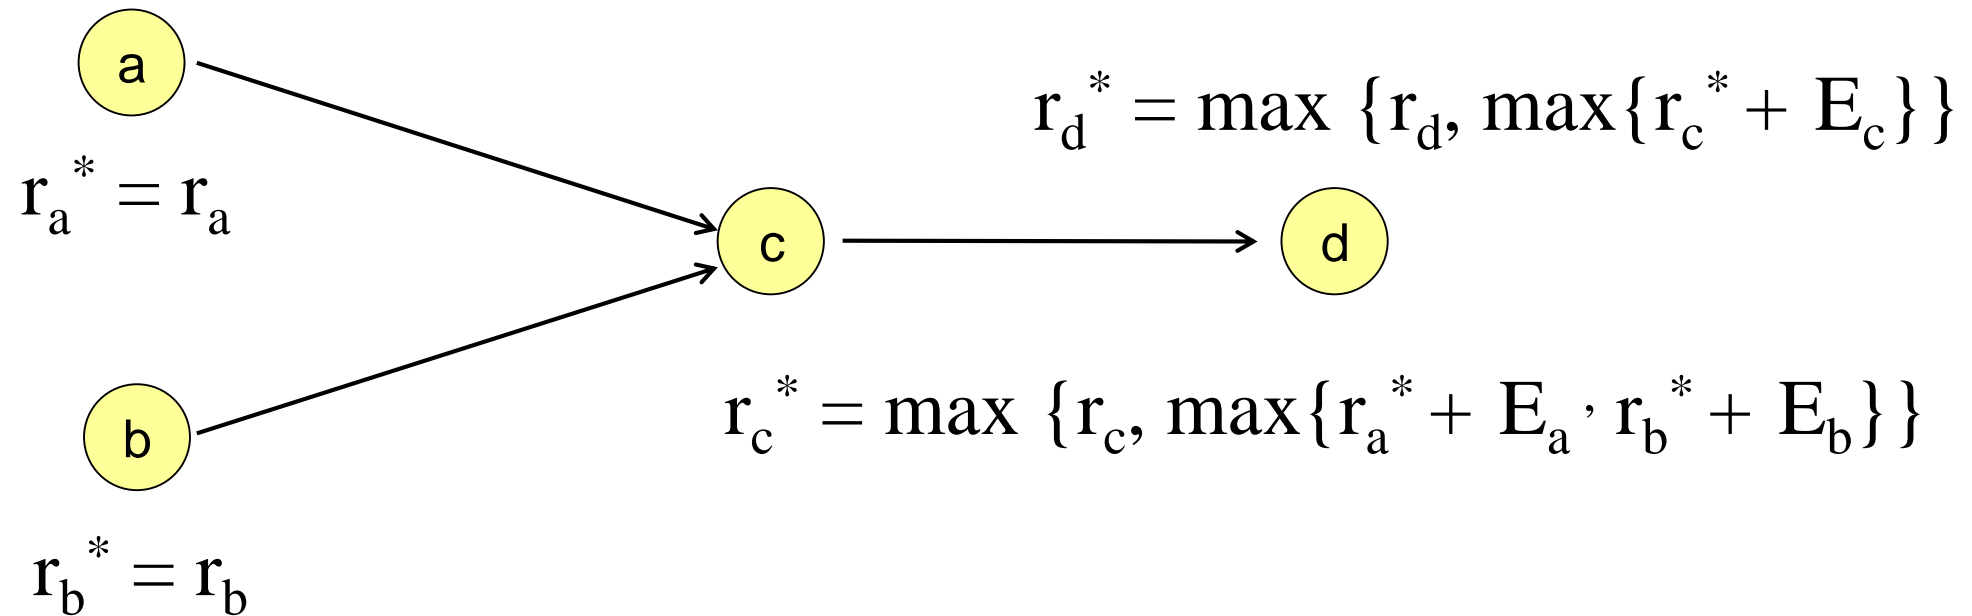
Example: EDF*



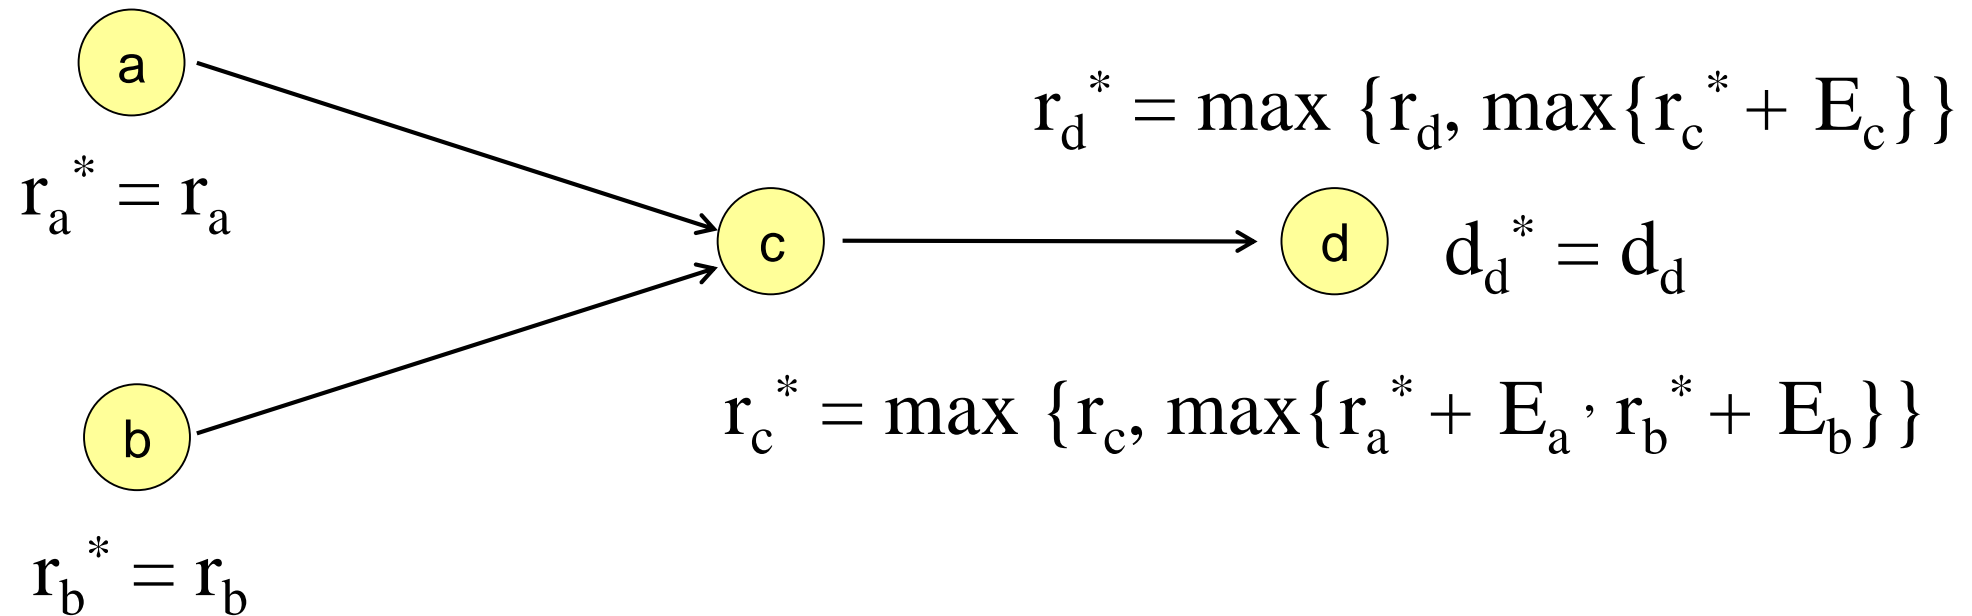
Example: EDF*



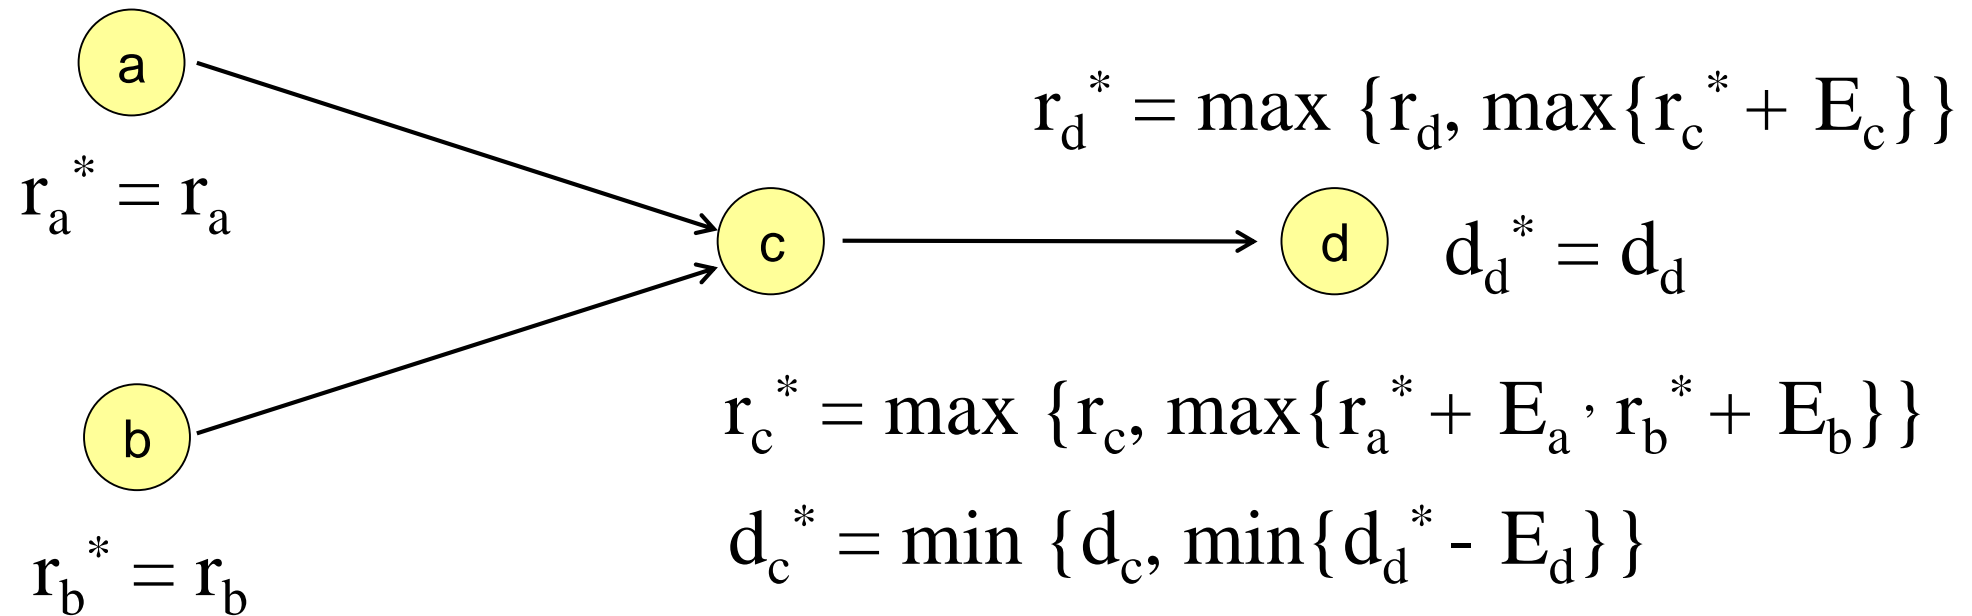
Example: EDF*



Example: EDF*




Example: EDF*



Example: EDF*

$$d_a^* = \min \{d_a, d_c^* - E_c\}$$



$$r_c^* = \max \{r_c, \max \{r_a^* + E_a, r_b^* + E_b\}\}$$

$$d_c^* = \min \{d_c, d_d^* - E_d\}$$

$$d_b^* = \min \{d_b, d_c^* - E_c\}$$

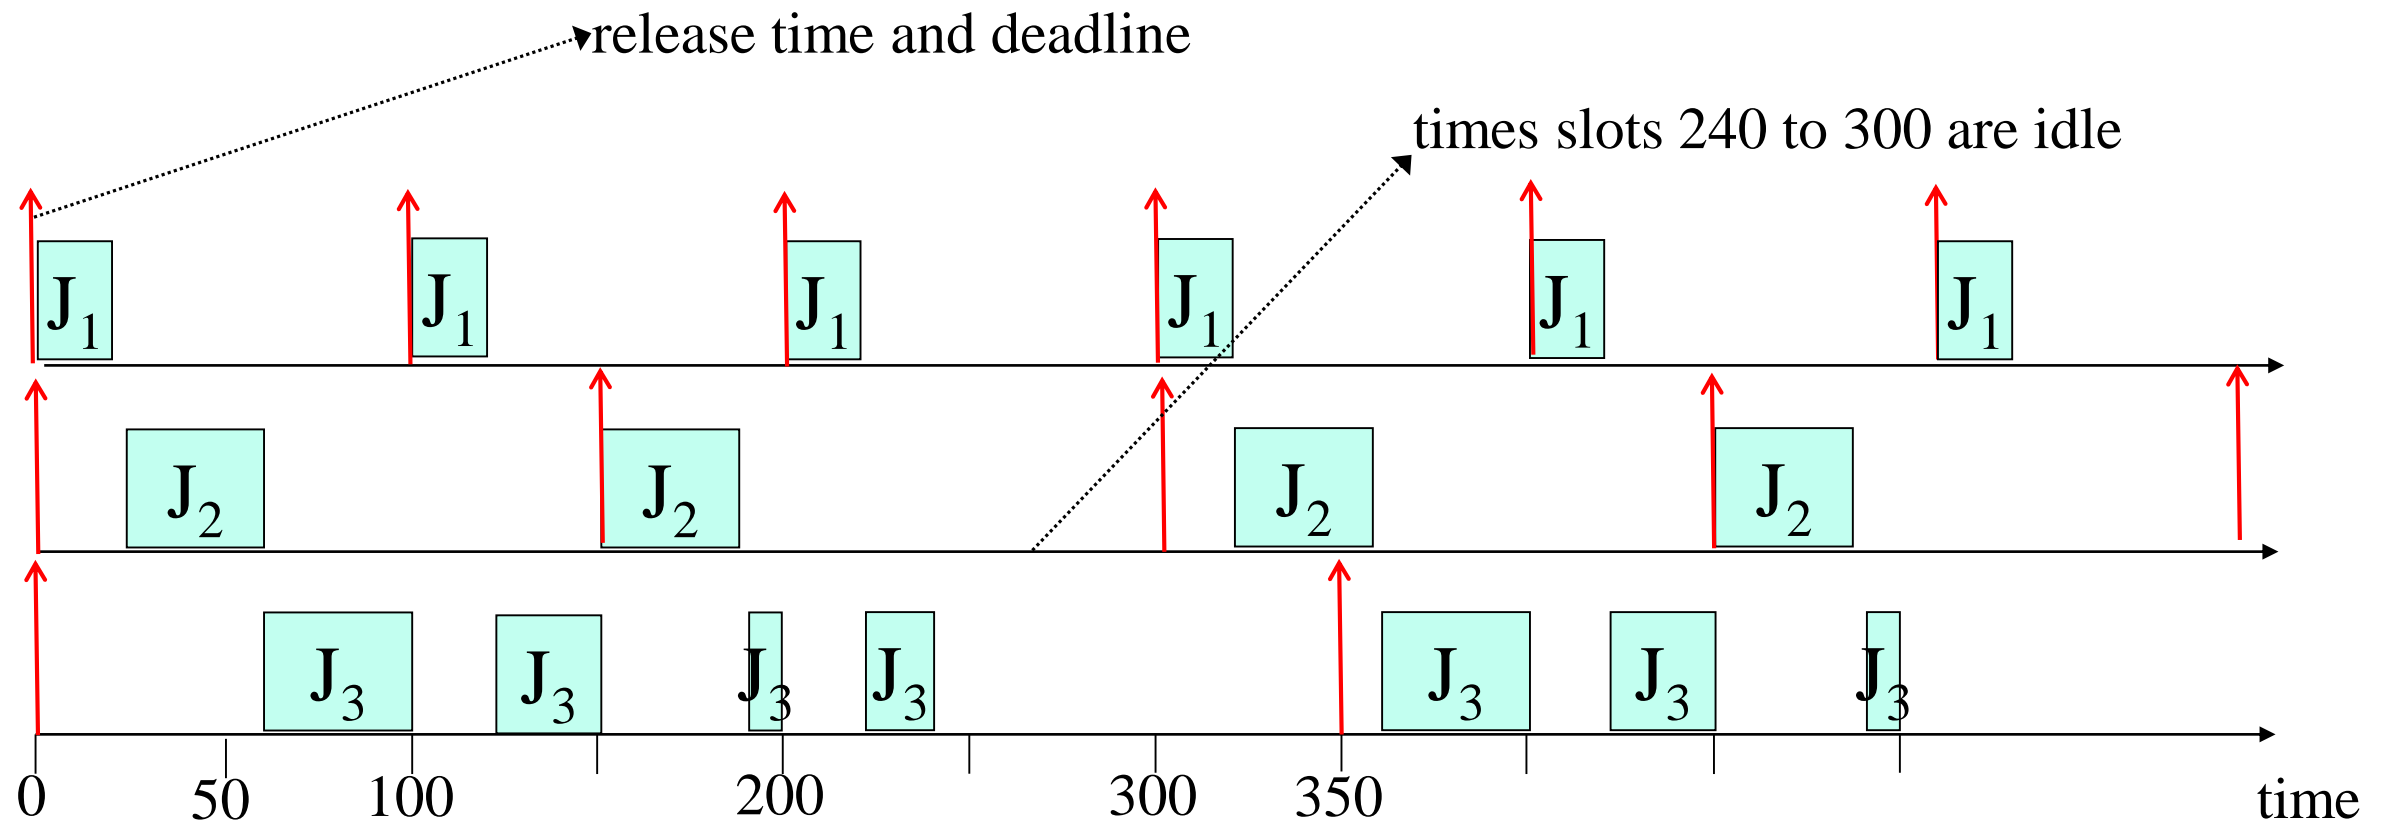
Rate Monotonic Scheduling (RM)

- For independent periodic tasks only
- Preemptive schedule
- Definition of priority:
Tasks with shorter periods have higher priorities
 - $\text{priority} = 1/\text{period}$ (this is called the rate)
- Priorities are static (they never change)
- Assumptions
 - $\text{deadline} = \text{period}$
 - tasks are always released at start of period

RM: Example 1

RM does produce feasible schedule!

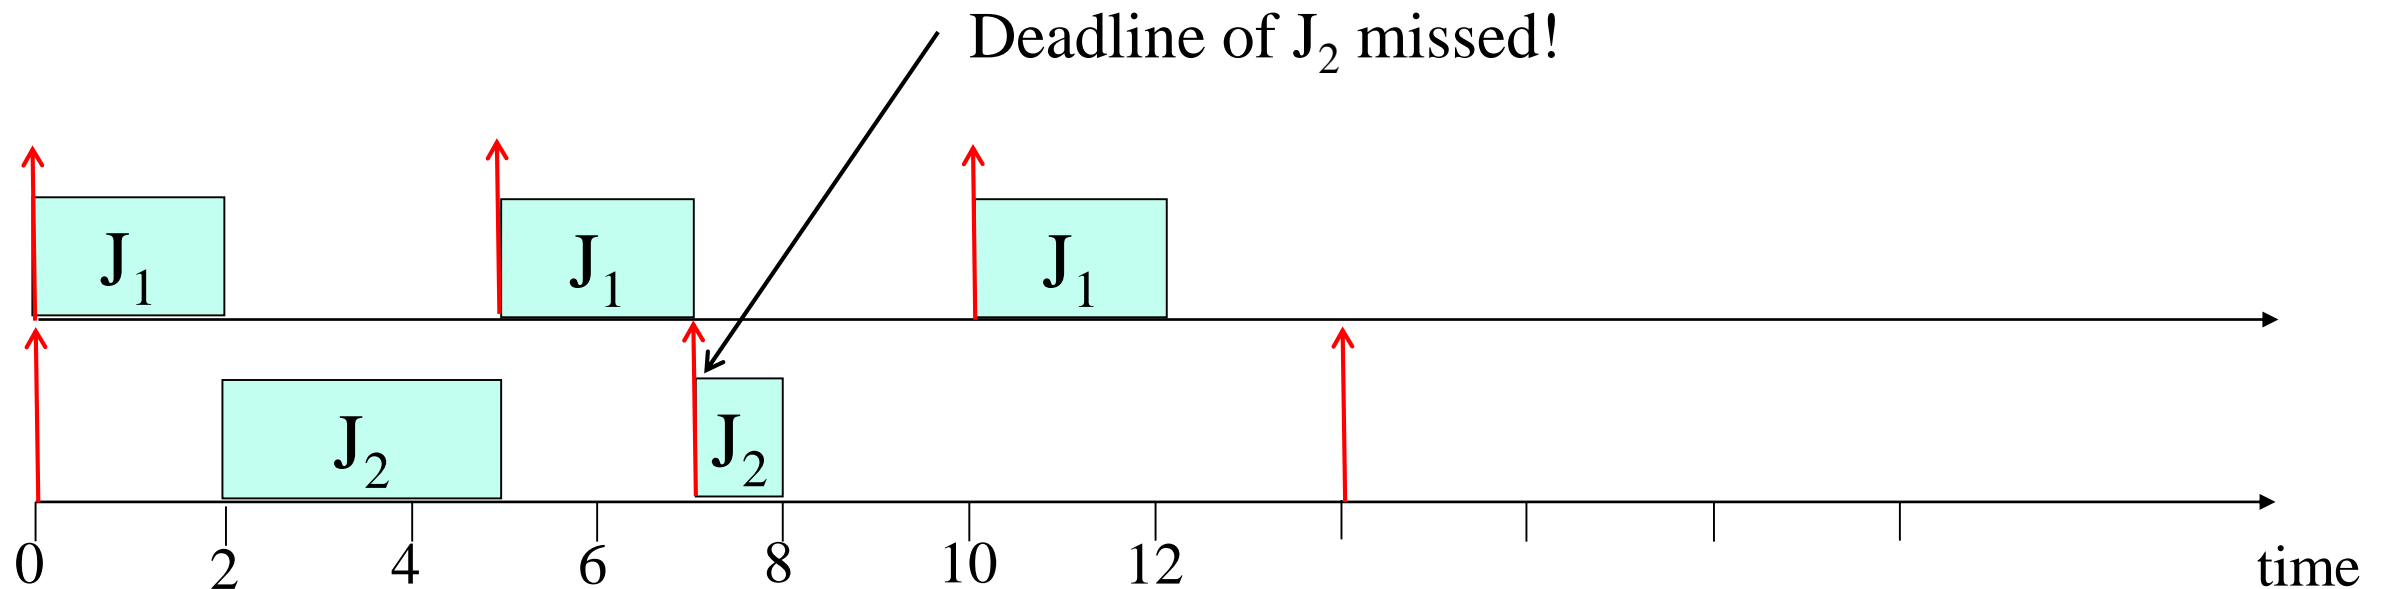
	J ₁	J ₂	J ₃
P _i	100	150	350
E _i	20	40	100



RM: Example 2

	J_1	J_2
P_i	5	7
E_i	2	4


RM does not always produce feasible schedule!



Example 2 with EDF

	J_1	J_2
P_i	5	7
E_i	2	4

EDF does produce a feasible schedule for this task set!



CPU Utilization

- CPU utilization U of a task is E/P
 - ratio of period P and execution time E
- CPU utilization a set of tasks is $U = \sum_i \frac{E_i}{P_i}$
 - $U = 0.752$ for first example
 - $U = 0.971$ for second example
- CPU utilization is a measure on how busy processor could be during shortest repeating cycle
 - $U > 1$ (overload)
 - some task will fail to meet deadline no matter what algorithm is used!
 - $U \leq 1$
 - doesn't imply schedulable, it depends on scheduling algorithm
 - If $U = 1$ and CPU is kept busy, all deadlines will be met
- **Theorem**
EDF produces a feasible schedule if and only if $U \leq 1$

Schedulability test for RM

- $U < 1$ does not imply schedulability for RM
 - see previous example
- The schedulability test for RMS is: $U \leq n(2^{1/n} - 1)$
 - For $n = 2$ this requires $U \leq 0.828$
 - For $n = 3$ this requires $U \leq 0.7797$
 - As number of tasks approaches grows, U converges to 0.69
 - i.e. if $U \leq 69\%$ then RMS guarantees to meet all deadlines
- RM is optimal in the following sense:
 - If a task set is schedulable with any fixed-priority scheduling algorithm, it is also schedulable with RM
 - EDF has no fixed priorities, priority = $1/(\text{time until deadline})$

RM: Example 3

$$U = 1/3 + 1/5 + 1/6 + 2/10 = 0.899$$


Schedulability bound: 0.756

RM produces a valid schedule

Hence, schedulability test is only sufficient, not necessary!

Precise/exact schedulability test exists

	J ₁	J ₂	J ₃	J ₄
P _i	3	5	6	10
E _i	1	1	1	2



Pros and Cons of RM

- Pros
 - Simple to understand
 - Easy to implement, static priority assignment
 - Stable: though some of the lower priority tasks fail to meet deadlines, others may meet deadlines
- Cons
 - lower CPU utilization than EDF
 - Requires $D=P$
 - Only deals with independent tasks
- The cons except lower CPU utilization can be fixed
- RM is no longer optimal if deadlines are shorter than period


Deadline Monotonic Scheduling (DM)

- Task model: same as for RM but deadlines can be shorter than period
- Definition of priority:
Tasks with shorter deadlines have higher priorities
- If deadlines are equal to periods then $DM = RM$

DM: Example 1

Time slot 17 is idle

	J ₁	J ₂	J ₃	J ₄
P _i	4	5	6	11
E _i	1	1	2	1
d _i	3	5	4	10



Deadline Monotonic Scheduling (DM)

- Principle
 - Scheduler keeps a list of released tasks, sorted by priorities
 - DM schedules first task on list
 - If a new task has higher priority than current task, then current task is preempted
- Schedulability test for DM
 - If $U = \sum_i \frac{E_i}{D_i} \leq n (2^{\frac{1}{n}} - 1)$ then task set schedule by DM
- Test is only sufficient
- $U = 0.874$ for previous example
- Bound for $n = 4$ is 0.756
- Is it schedulable?



Real-time Operating Systems

Real-time Operating Systems (RTOS)

- RTOS
Operating system designed to meet strict deadlines
- Methods employed
 - Drop or reduce certain tasks when they cannot be executed within time constraints (load shedding)
 - Monitor input consistently and in a timely manner
 - Keep track of how much of each resource (CPU time, RAM, communications bandwidth, etc.) might possibly be used in worst-case by currently-running task, and refuse to accept a new task unless it fits in remaining unallocated resources

Simple Variant of RTOS

- Periodic execution
 - Provides a minimal time service: Scheduled clock pulse with fixed period
 - No preemption
 - Allows implementation of a static cyclic schedule, provided:
 - All tasks can be scheduled in a frame-based manner according to their WCET
 - All interactions with hardware to be done on a polled basis, i.e. no blocking

Simple Variant of RTOS

- Common approach:

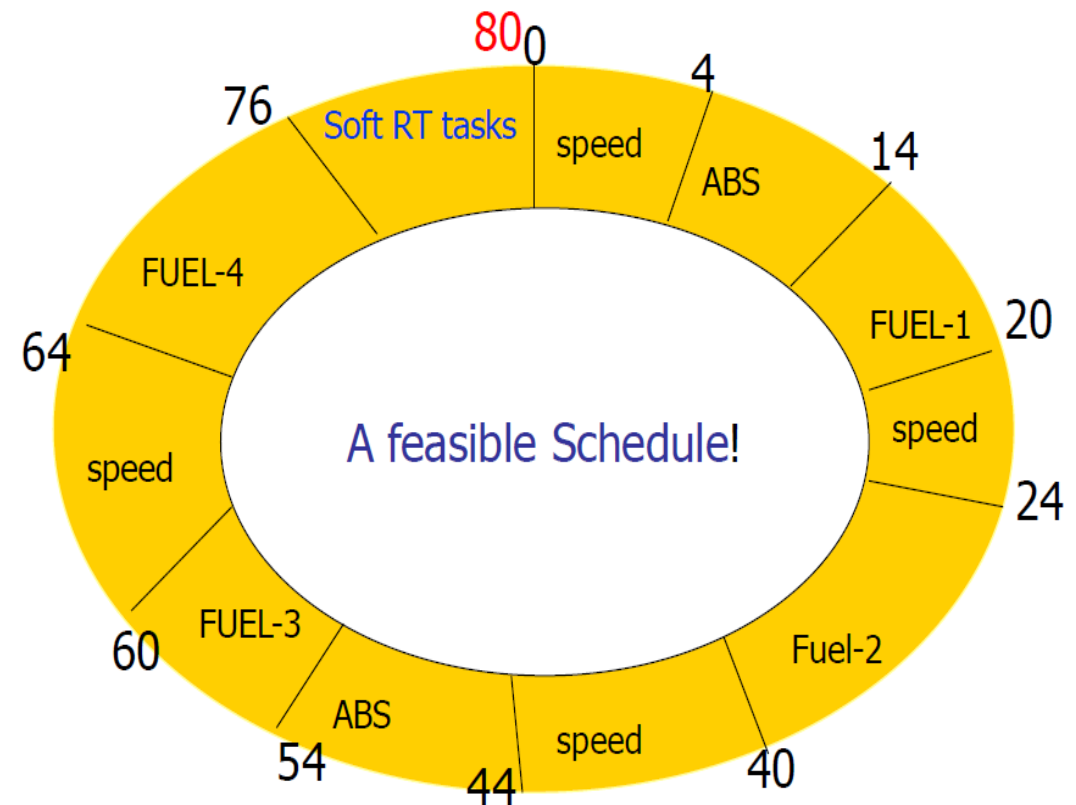
```
setup timer
c = 0;
while (1) {
    suspend until timer expires
    c++;
    do tasks due every cycle
    if ((c % 2) == 0)
        do tasks due every 2nd cycle
    if ((c % 3) == 0)
        do tasks due every 3rd cycle
    ...
}
```

Advanced Task scheduling

- Many real-time embedded systems are more complex, they need a sophisticated OS with priority scheduling
- Compiler generates execution plan (schedule) based on information about WCET
 - Non-preemptive schedules
 - Only determines order of tasks, tasks always finish execution
 - Preemptive schedules
 - Contains information about order and how long a task executes before preemption to allow next task to execute

Example 2: Car Control

- Activities of a car control system
 - Speed measurement: $E = 4$ ms, $P = 20$ ms
 - ABS control: $E = 10$ ms, $P = 40$ ms
 - Fuel injection: $E = 40$ ms, $P = 80$ ms
 - Other soft deadlines
 - air condition etc.
 - Schedule produced by EDF



Existing RTOS: Windows CE

- Built specifically for embedded systems and appliance market
- Scalable real-time 32-bit platform, supports Windows API
- Preemptive priority scheduling with 256 priority levels per process/task
- Kernel is 400 Kbytes
- Development tool: Visual Studio
- Supports Intel x86 and compatibles, MIPS, and ARM processors

Existing RTOS: QNX

- Real-time microkernel surrounded by resource managers
 - Provides POSIX/UNIX compatibility
 - Microkernels typically support only most basic services
 - Optional resource managers allow scalability from small ROM-based systems to huge multiprocessor systems
- Preemptive task scheduling using FIFO, round-robin, adaptive, or priorities
- Microkernel < 10 Kbytes and complies with POSIX real-time standard
- Supports PowerPC, x86 family, MIPS, ARM, StrongARM and XScale CPUs

- Alternatives: RTLinux, LynxOS, VxWorks, eCos,...

SES

Chapter 10: Real-Time Embedded Systems

Prof. Dr.-Ing. Bernd-Christian Renner

